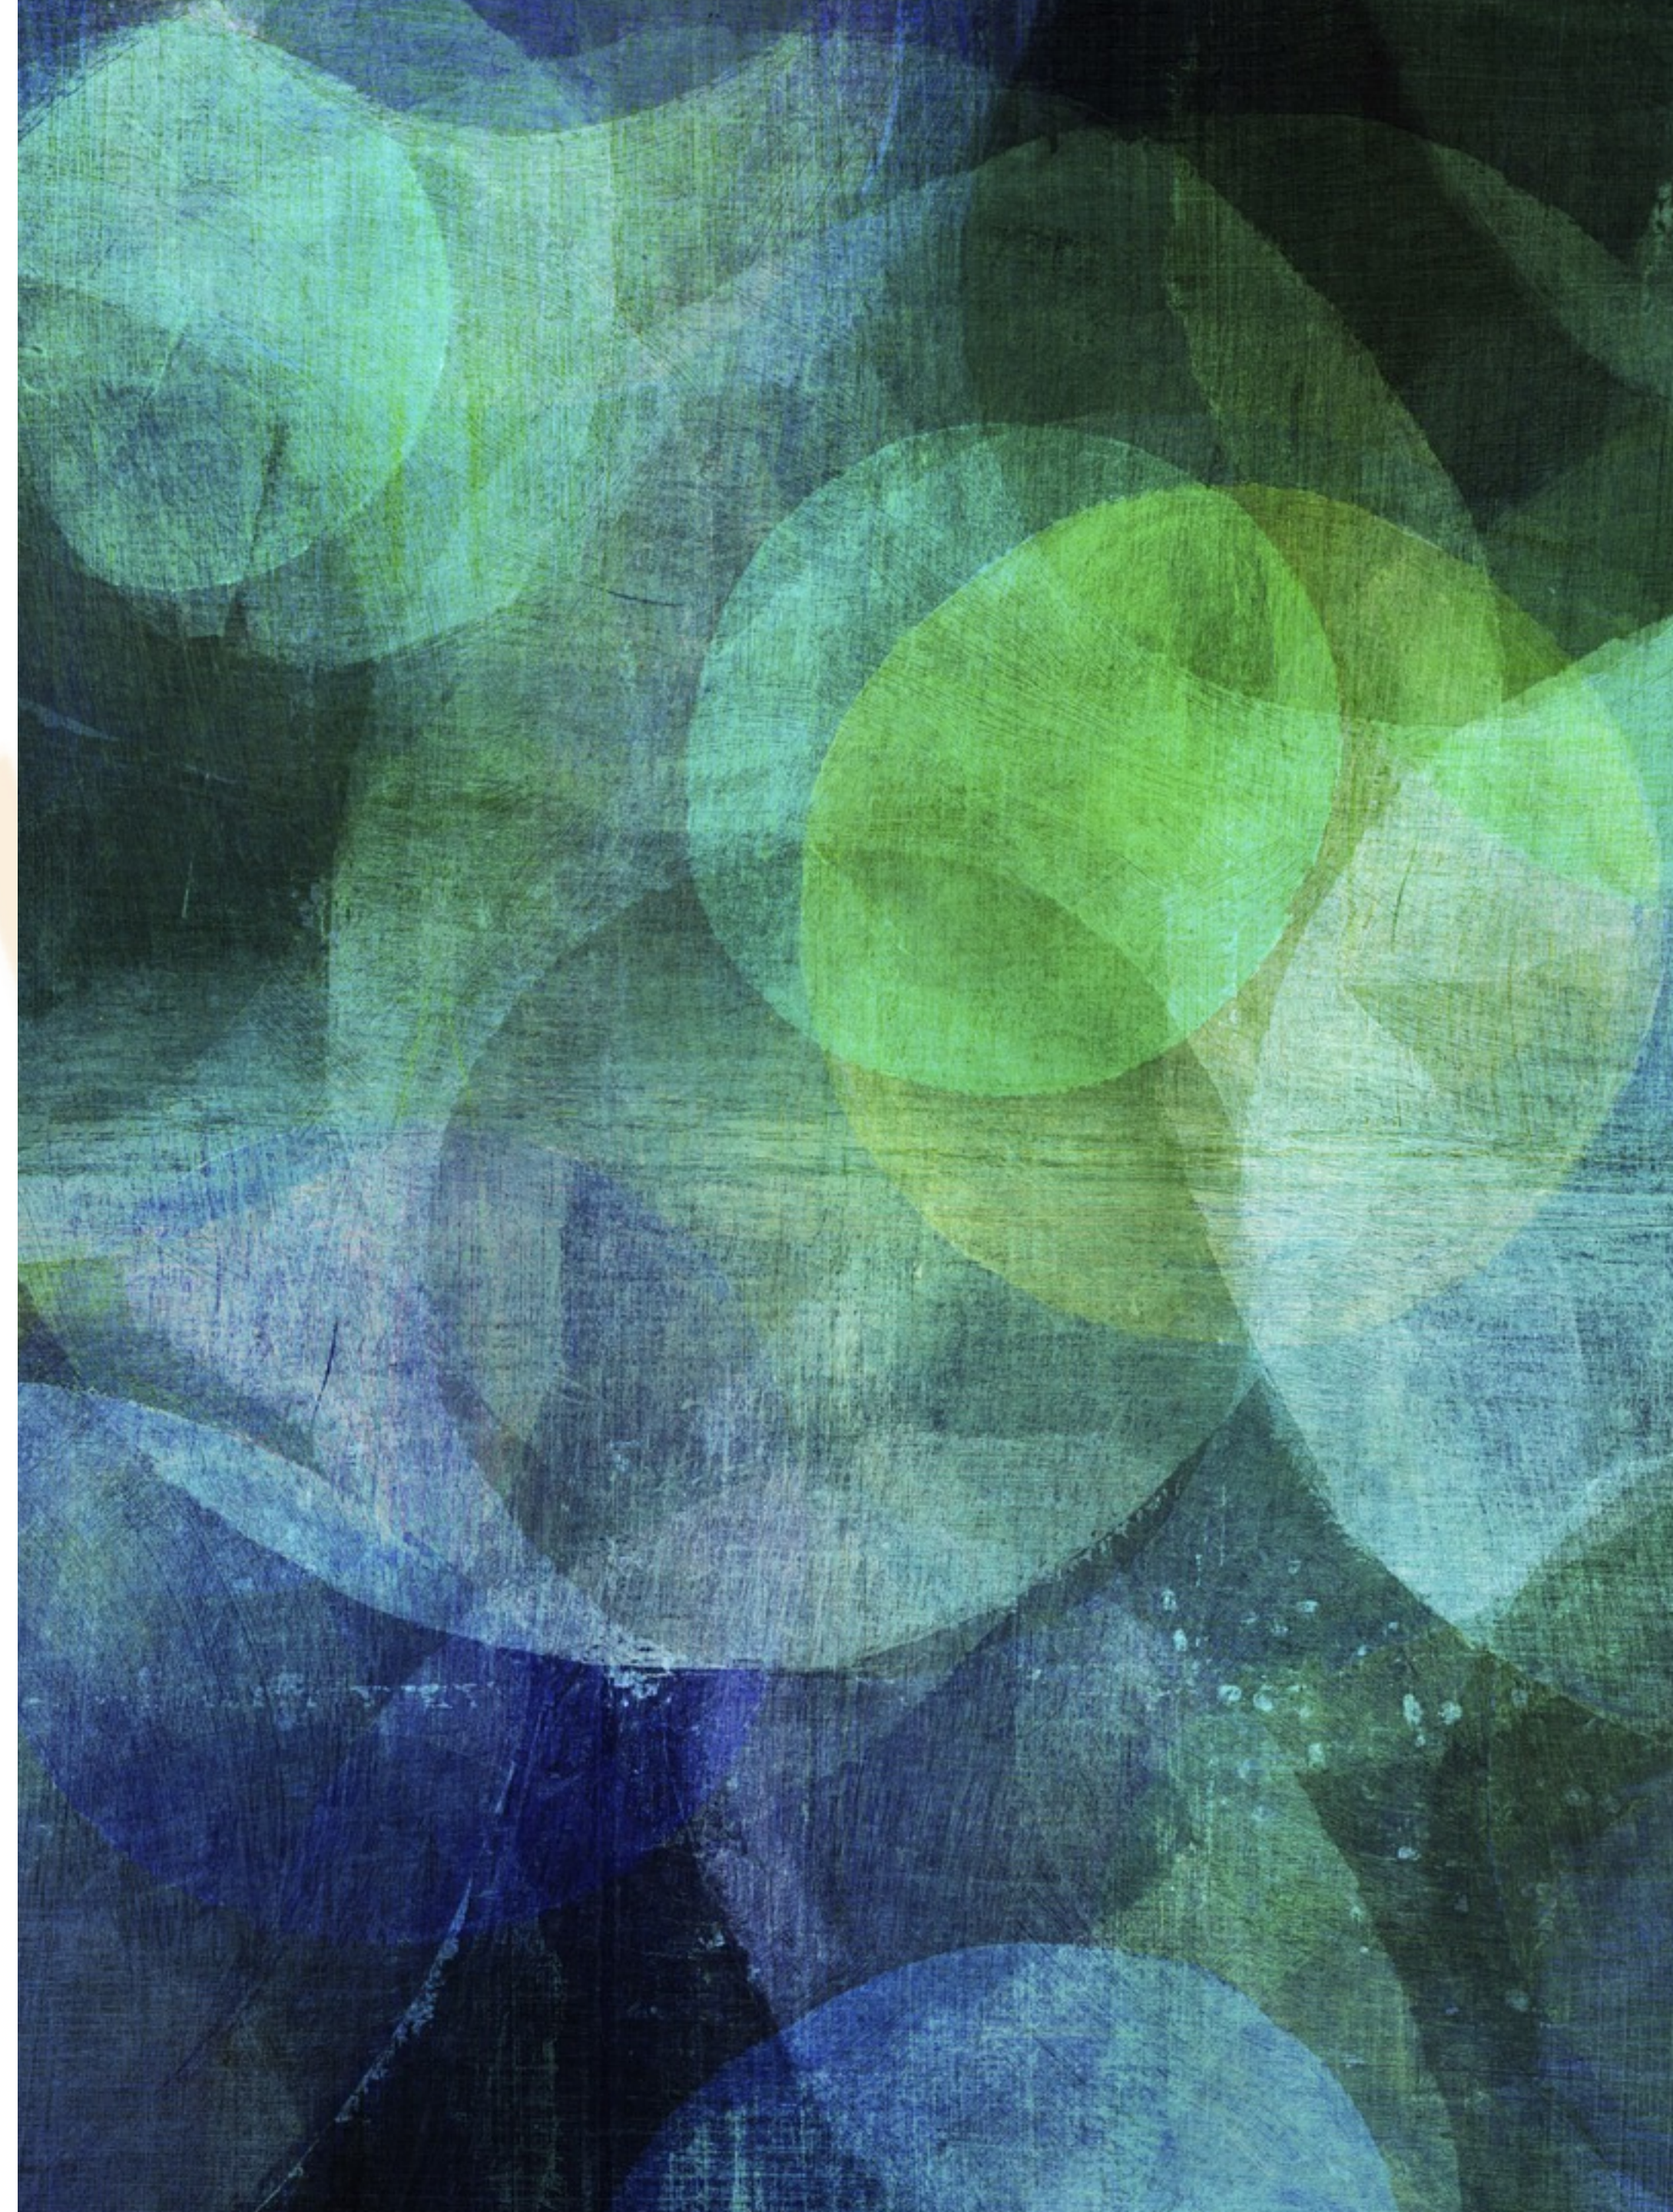


AWS DATA PROCESSING INFRASTRUCTURE 4A

Nan Dun
nan.dun@acm.org



COPYRIGHT POLICY 版权声明

All content included on the Site or third-party platforms as part of the class, such as text, graphics, logos, button icons, images, audio clips, video clips, live streams, digital downloads, data compilations, and software, is the property of BitTiger or its content suppliers and protected by copyright laws.

Any attempt to redistribute or resell BitTiger content will result in the appropriate legal action being taken.

We thank you in advance for respecting our copyrighted content.

For more info see <https://www.bittiger.io/termsfuse> and <https://www.bittiger.io/termservice>

所有太阁官方网站以及在第三方平台课程中所产生的课程内容，如文本，图形，徽标，按钮图标，图像，音频剪辑，视频剪辑，直播流，数字下载，数据编辑和软件均属于太阁所有并受版权法保护。

对于任何尝试散播或转售BitTiger的所属资料的行为，太阁将采取适当的法律行动。

我们非常感谢您尊重我们的版权内容。

有关详情，请参阅

<https://www.bittiger.io/termsfuse>

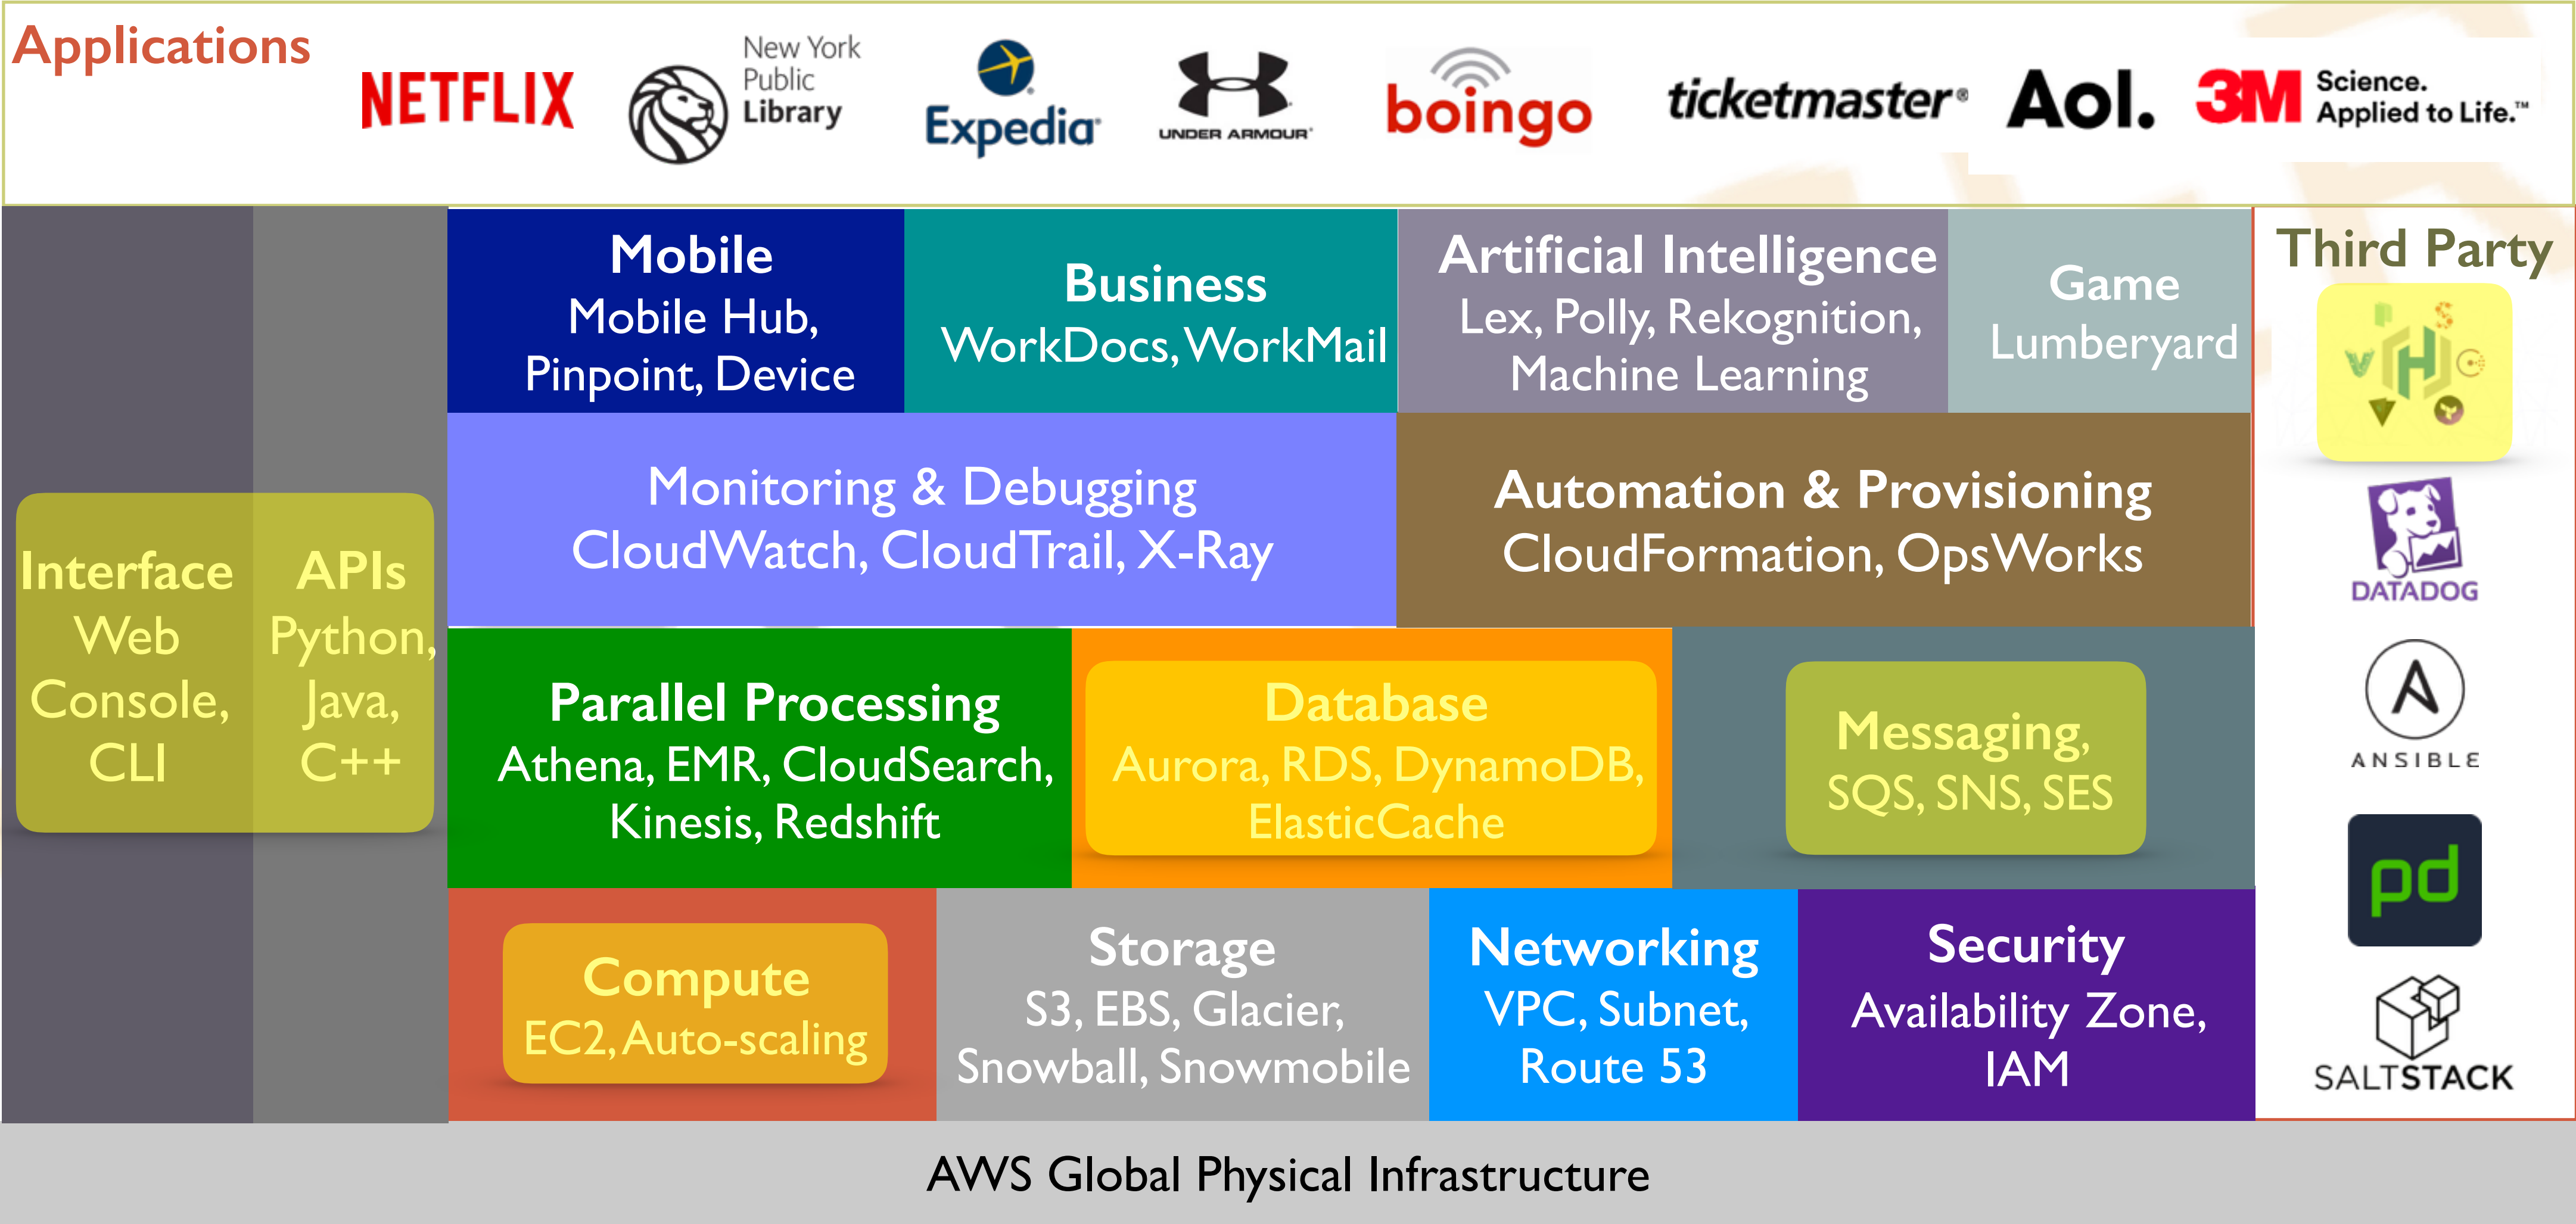
<https://www.bittiger.io/termservice>



DISCLAIMER

- I. “All data, information, and opinions expressed in this presentation is for informational purposes only. I do not guarantee the accuracy or reliability of the information provided herein. This is a personal presentation. The opinions expressed here represent my own and not those of my employer.”*
- II. “The copyright of photos, icons, charts, trademarks presented here belong to their authors.”*
- III. “I could be wrong.”*

TODAY'S TOPIC



OUTLINE

- Simple Queue Service (SQS)
- DynamoDB
- Docker and Container
- Elastic Container Service (ECS)
- NYC Taxi Trip Explorer Project





BITTIGER

SQS

SQS

- A reliable, and highly-scalable queue (message send/recv) service between applications
- For what purpose
 - Decoupling (one most import design principle for large-scale distributed systems)
 - Concurrency
 - Batch processing
 - Buffering
 - Coordinating
 - ...
- Core concepts: Queue, message, send, receive, delete

QUEUE OPERATIONS

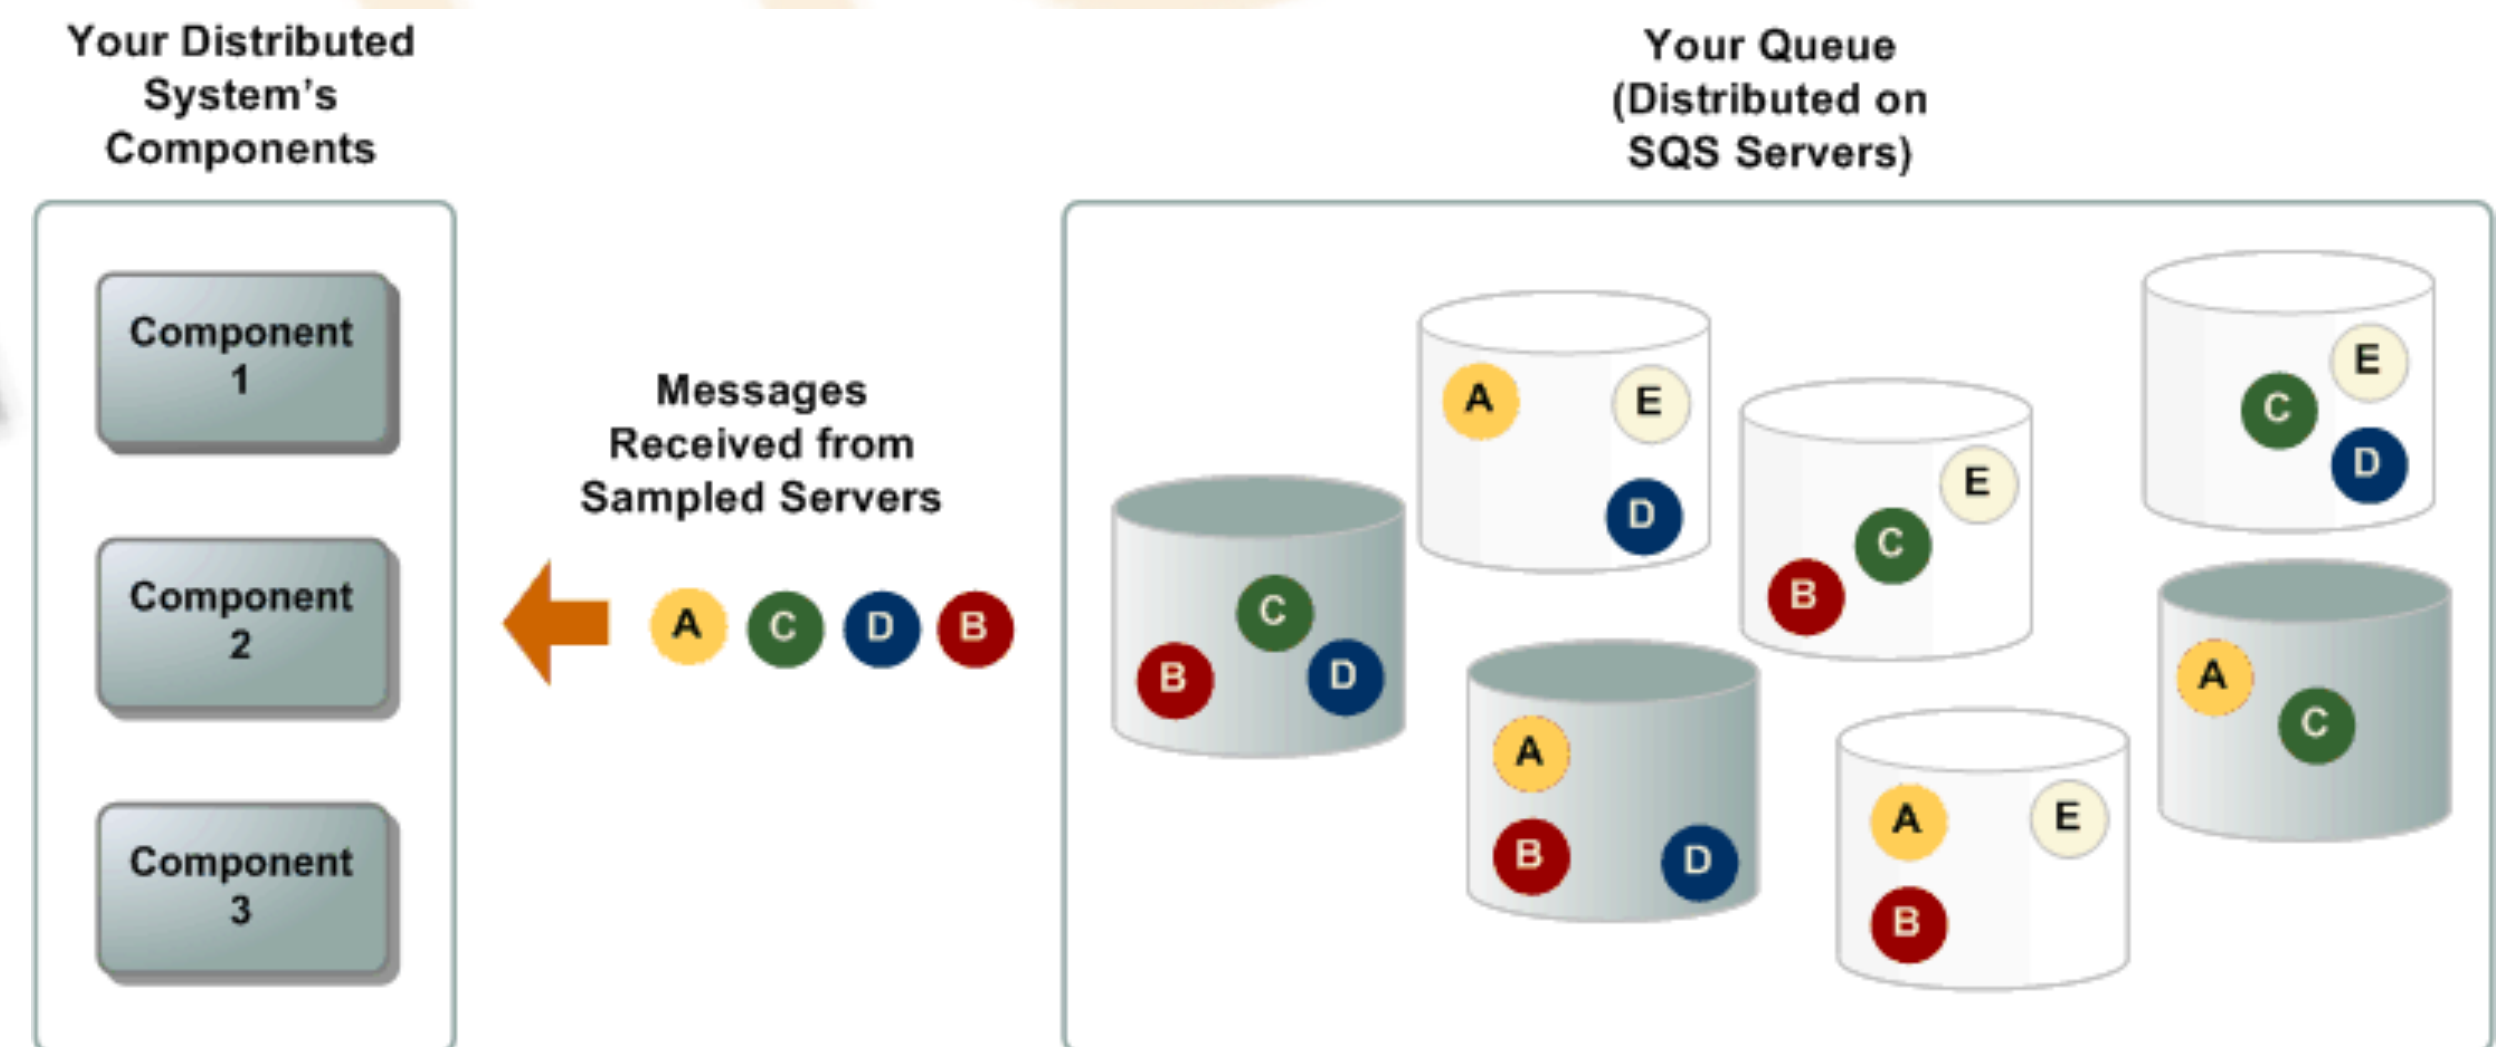
- Create a queue
- List all queues
- Add permission to a queue
 - IAM policies
- Purging a queue
- Delete a queue
- Subscribing a queue to a SNS topic
 - Store SNS notification to you queue

MESSAGE OPERATIONS

- Send
 - One message up to 256 KB, encoded as a string.
 - Messages can be sent in bulks of up to 10 (but the total size is capped at 256 KB).
- Receive
 - Up to 10 messages can be received in bulk, if available in the queue.
- Long polling
 - The request will wait up to 20 seconds for messages, if none are available initially
- Delete

STANDARD QUEUES

- Message order
 - best effort
- At-least-Once delivery
 - Receive or delete may fail, receive same message again
- Receive message by short polling
 - Sampling and deliver
 - Request may not be fulfilled
- Unlimited transactions



FIFO QUEUES

- Strict order of messages
 - Strict receive order as sent order
- Exactly-once processing
 - By providing a duplication ID
 - Or content-based
- Limited to 300 transactions per second
- Available only in us-west-2(oregon) and us-east-1(ohio)


QUEUE ATTRIBUTES

- Approximate number of messages
- Approximate number of message delayed
- Approximate number of message not visible



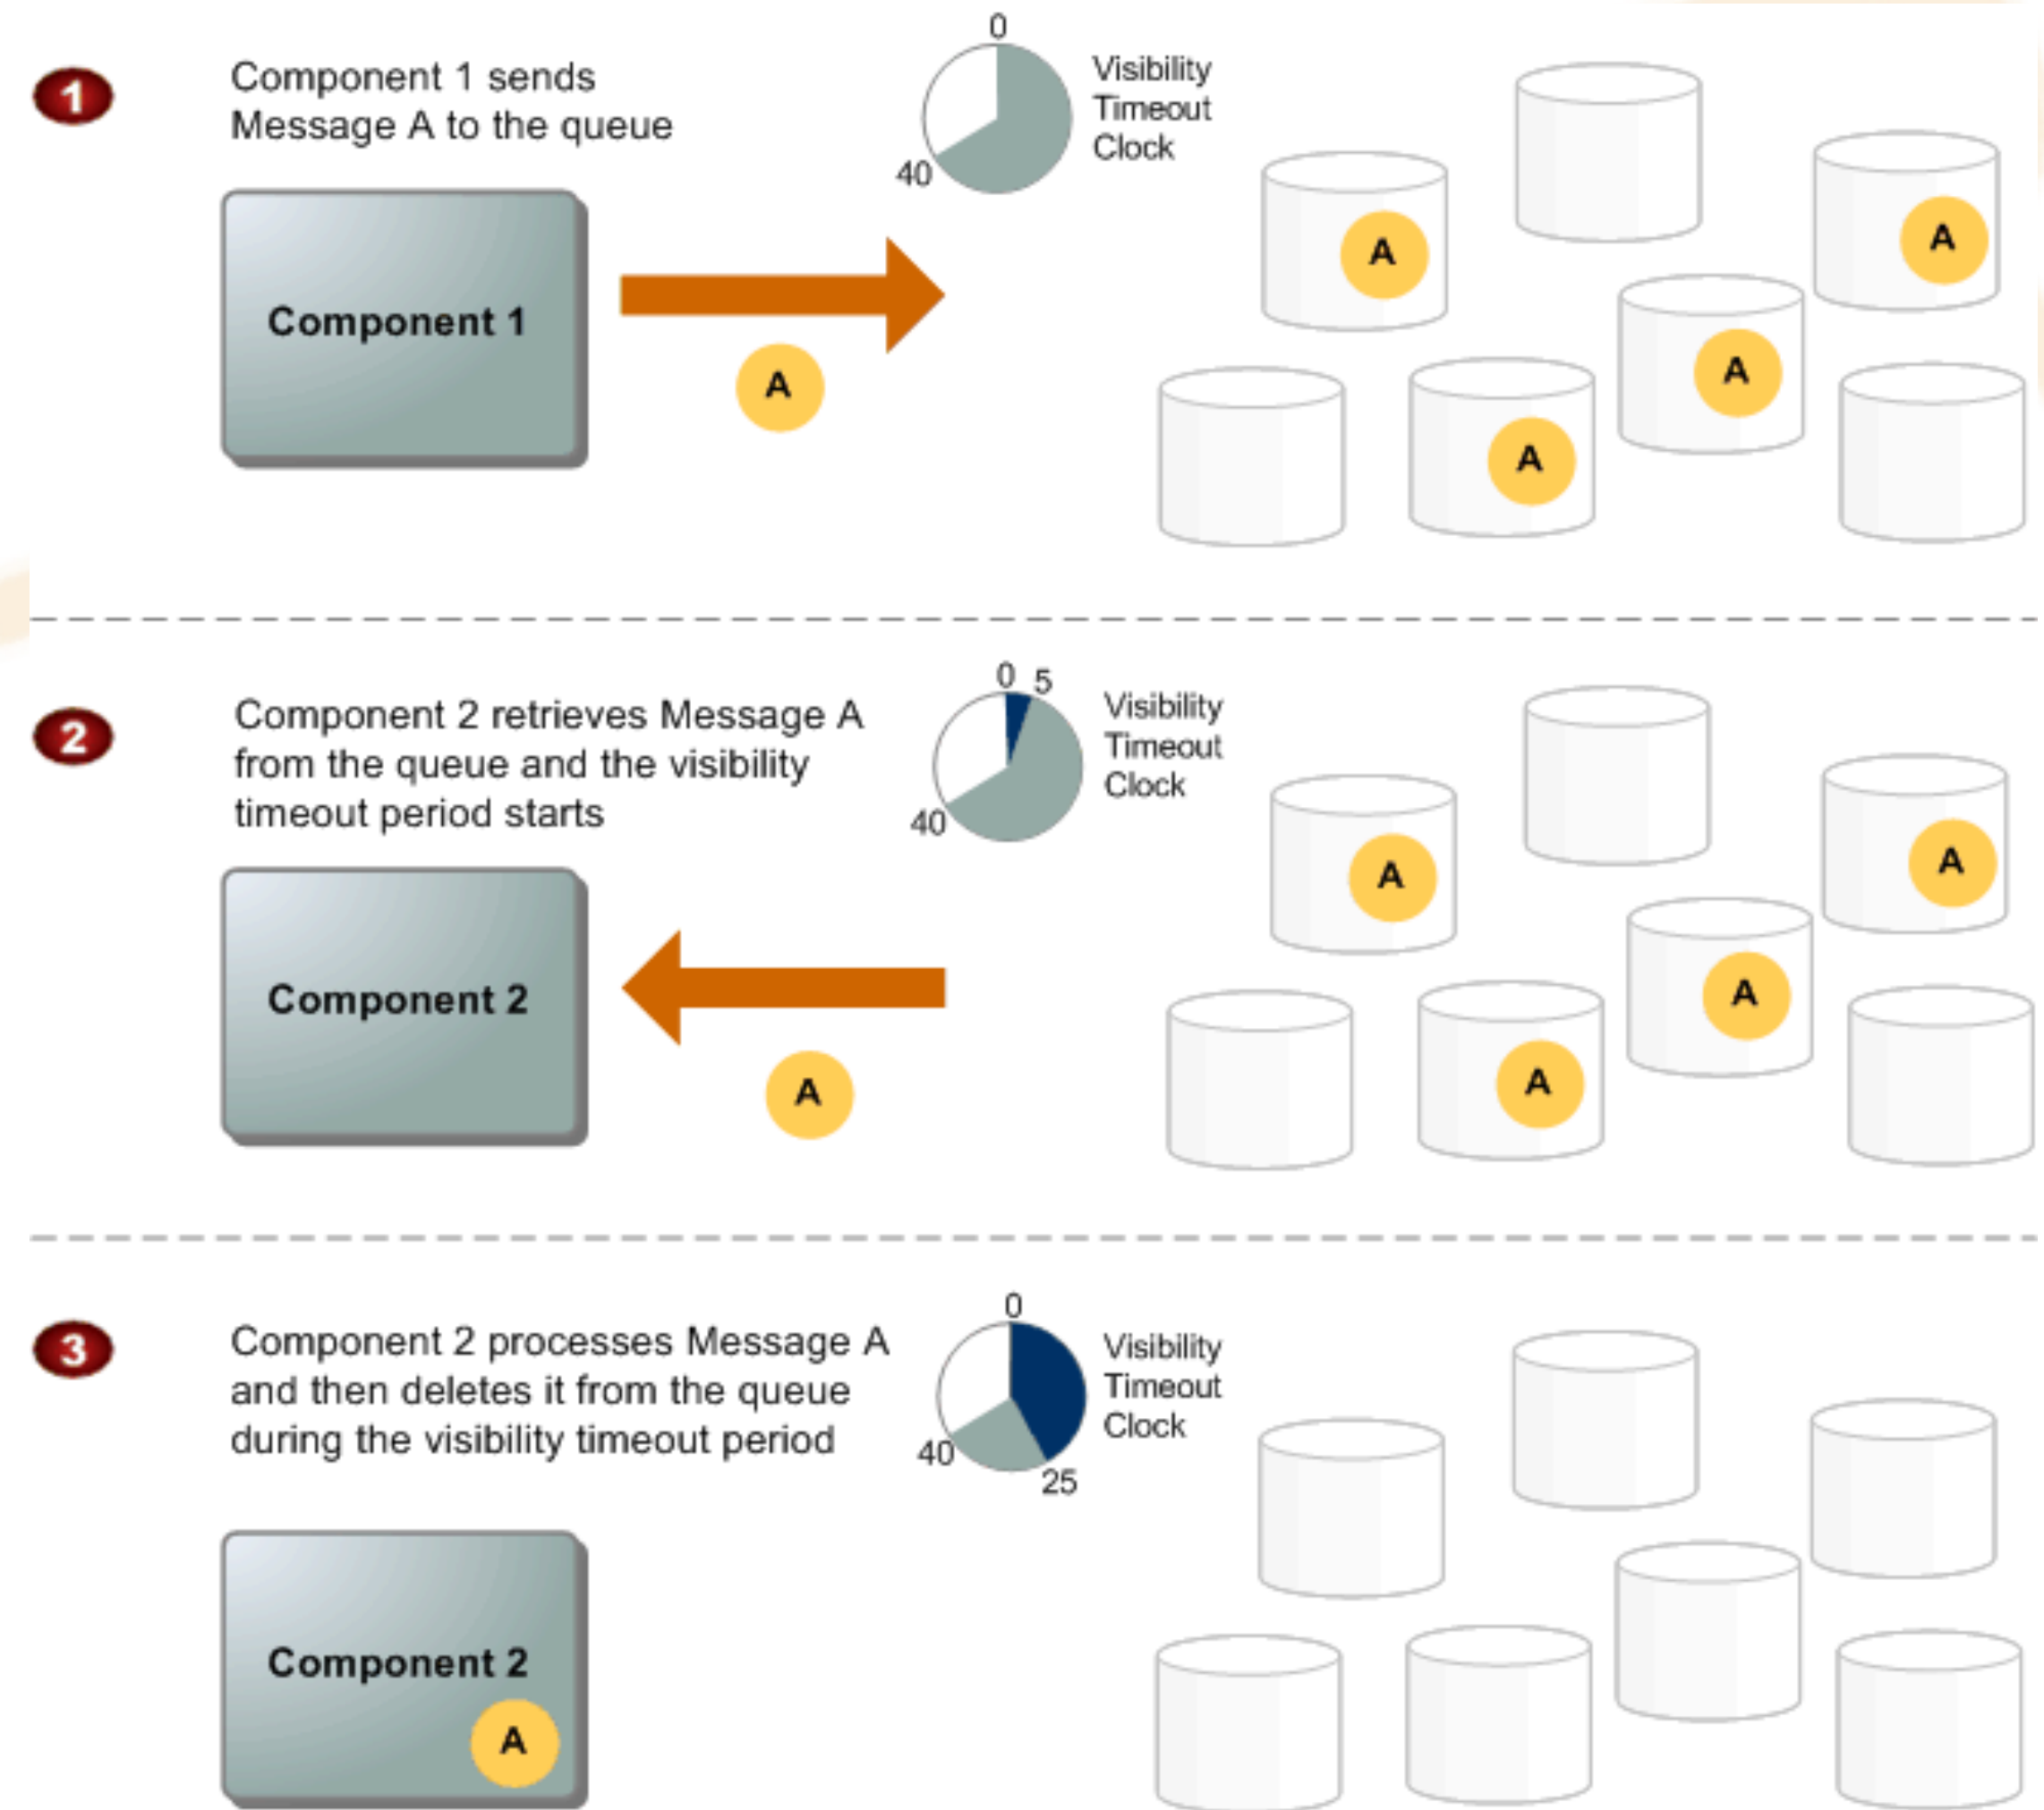
BITTIGER

VISIBILITY TIMEOUT



MESSAGE LIFECYCLE

- Send and distribute
- Receive and is hidden for visibility timeout
- Delete a message or receive again after visibility timeout



DEAD LETTER QUEUES

- Receive messages for
 - Message that is sent to a queue that does not exist
 - Target queue is full
 - Message length limit exceeded
 - Any other failures



MESSAGE ATTRIBUTES

- User-defined
 - Name
 - Type
 - Value
- Up to 10 attributes
- Used to help process the message



LONG POLLING

- Polling for longer than 20 seconds
- To
 - Reduce the number of empty responses
 - Eliminate false empty message by querying all servers
 - Returns as soon as message become available



MESSAGE TIMER


- Message's internal visibility timeout, not queue default
- Use to
 - Delay specific messages
 - Implement task priority



BITTIGER

DELAY QUEUES AND MESSAGE TIMERS

- Set a initial invisibility timeout
 - For default queue
 - or, message
- Message should wait for timeout to be visible
- 120,000 inflight message to reach OverLimit error
- For
 - Buffering




LARGE MESSAGES?


- 256 KB < message size < 2GB
- Data stored in S3, use message as a pointer/reference
- Operations
 - Select to store all on S3 or only when size > 256KB
 - Send a reference of message on S3
 - Receive a message on S3
 - Delete
- Java only with extended client library
 - <https://github.com/aws-labs/amazon-sqs-java-extended-client-lib>

CAPACITY

- EC2 m2.xlarge
 - 100K messages/minute
 - 1M messages/minute by 10 instances
- Depends on your applications



Source: <http://www.warski.org/blog/2014/06/benchmarking-sqs/>



DYNAMODB



BITTIGER

NOSQL VS. SQL

- NoSQL examples
 - Column
 - Key-value store
 - Document
 - Graph
 - ...



BITTIGER

CONCEPTS

- Tables
- Items
- Attributes
- Primary key
 - Partition key
 - Partition key + sorted key
- Secondary indexes
 - Global secondary index
 - Partition key + sorted key different from primary
 - Local secondary index
 - Same partition key as default

People


```
{
  "PersonID": 101,
  "LastName": "Smith",
  "FirstName": "Fred",
  "Phone": "555-4321"
}

{
  "PersonID": 102,
  "LastName": "Jones",
  "FirstName": "Mary",
  "Address": {
    "Street": "123 Main",
    "City": "Anytown",
    "State": "OH",
    "ZIPCode": 12345
  }
}

{
  "PersonID": 103,
  "LastName": "Stephens",
  "FirstName": "Howard",
  "Address": {
    "Street": "123 Main",
    "City": "London",
    "PostalCode": "ER3 5K8"
  },
  "FavoriteColor": "Blue"
}
```


DB OPERATIONS

- Table
 - Create
 - Describe
 - List
 - Update
 - Delete
- Data
 - Create: String|number|binary|bool|null
 - Read
 - Get, BatchGet
 - Query
 - Scan
 - Update
 - Delete
 - Delete
 - BatchWrite: del 25 once



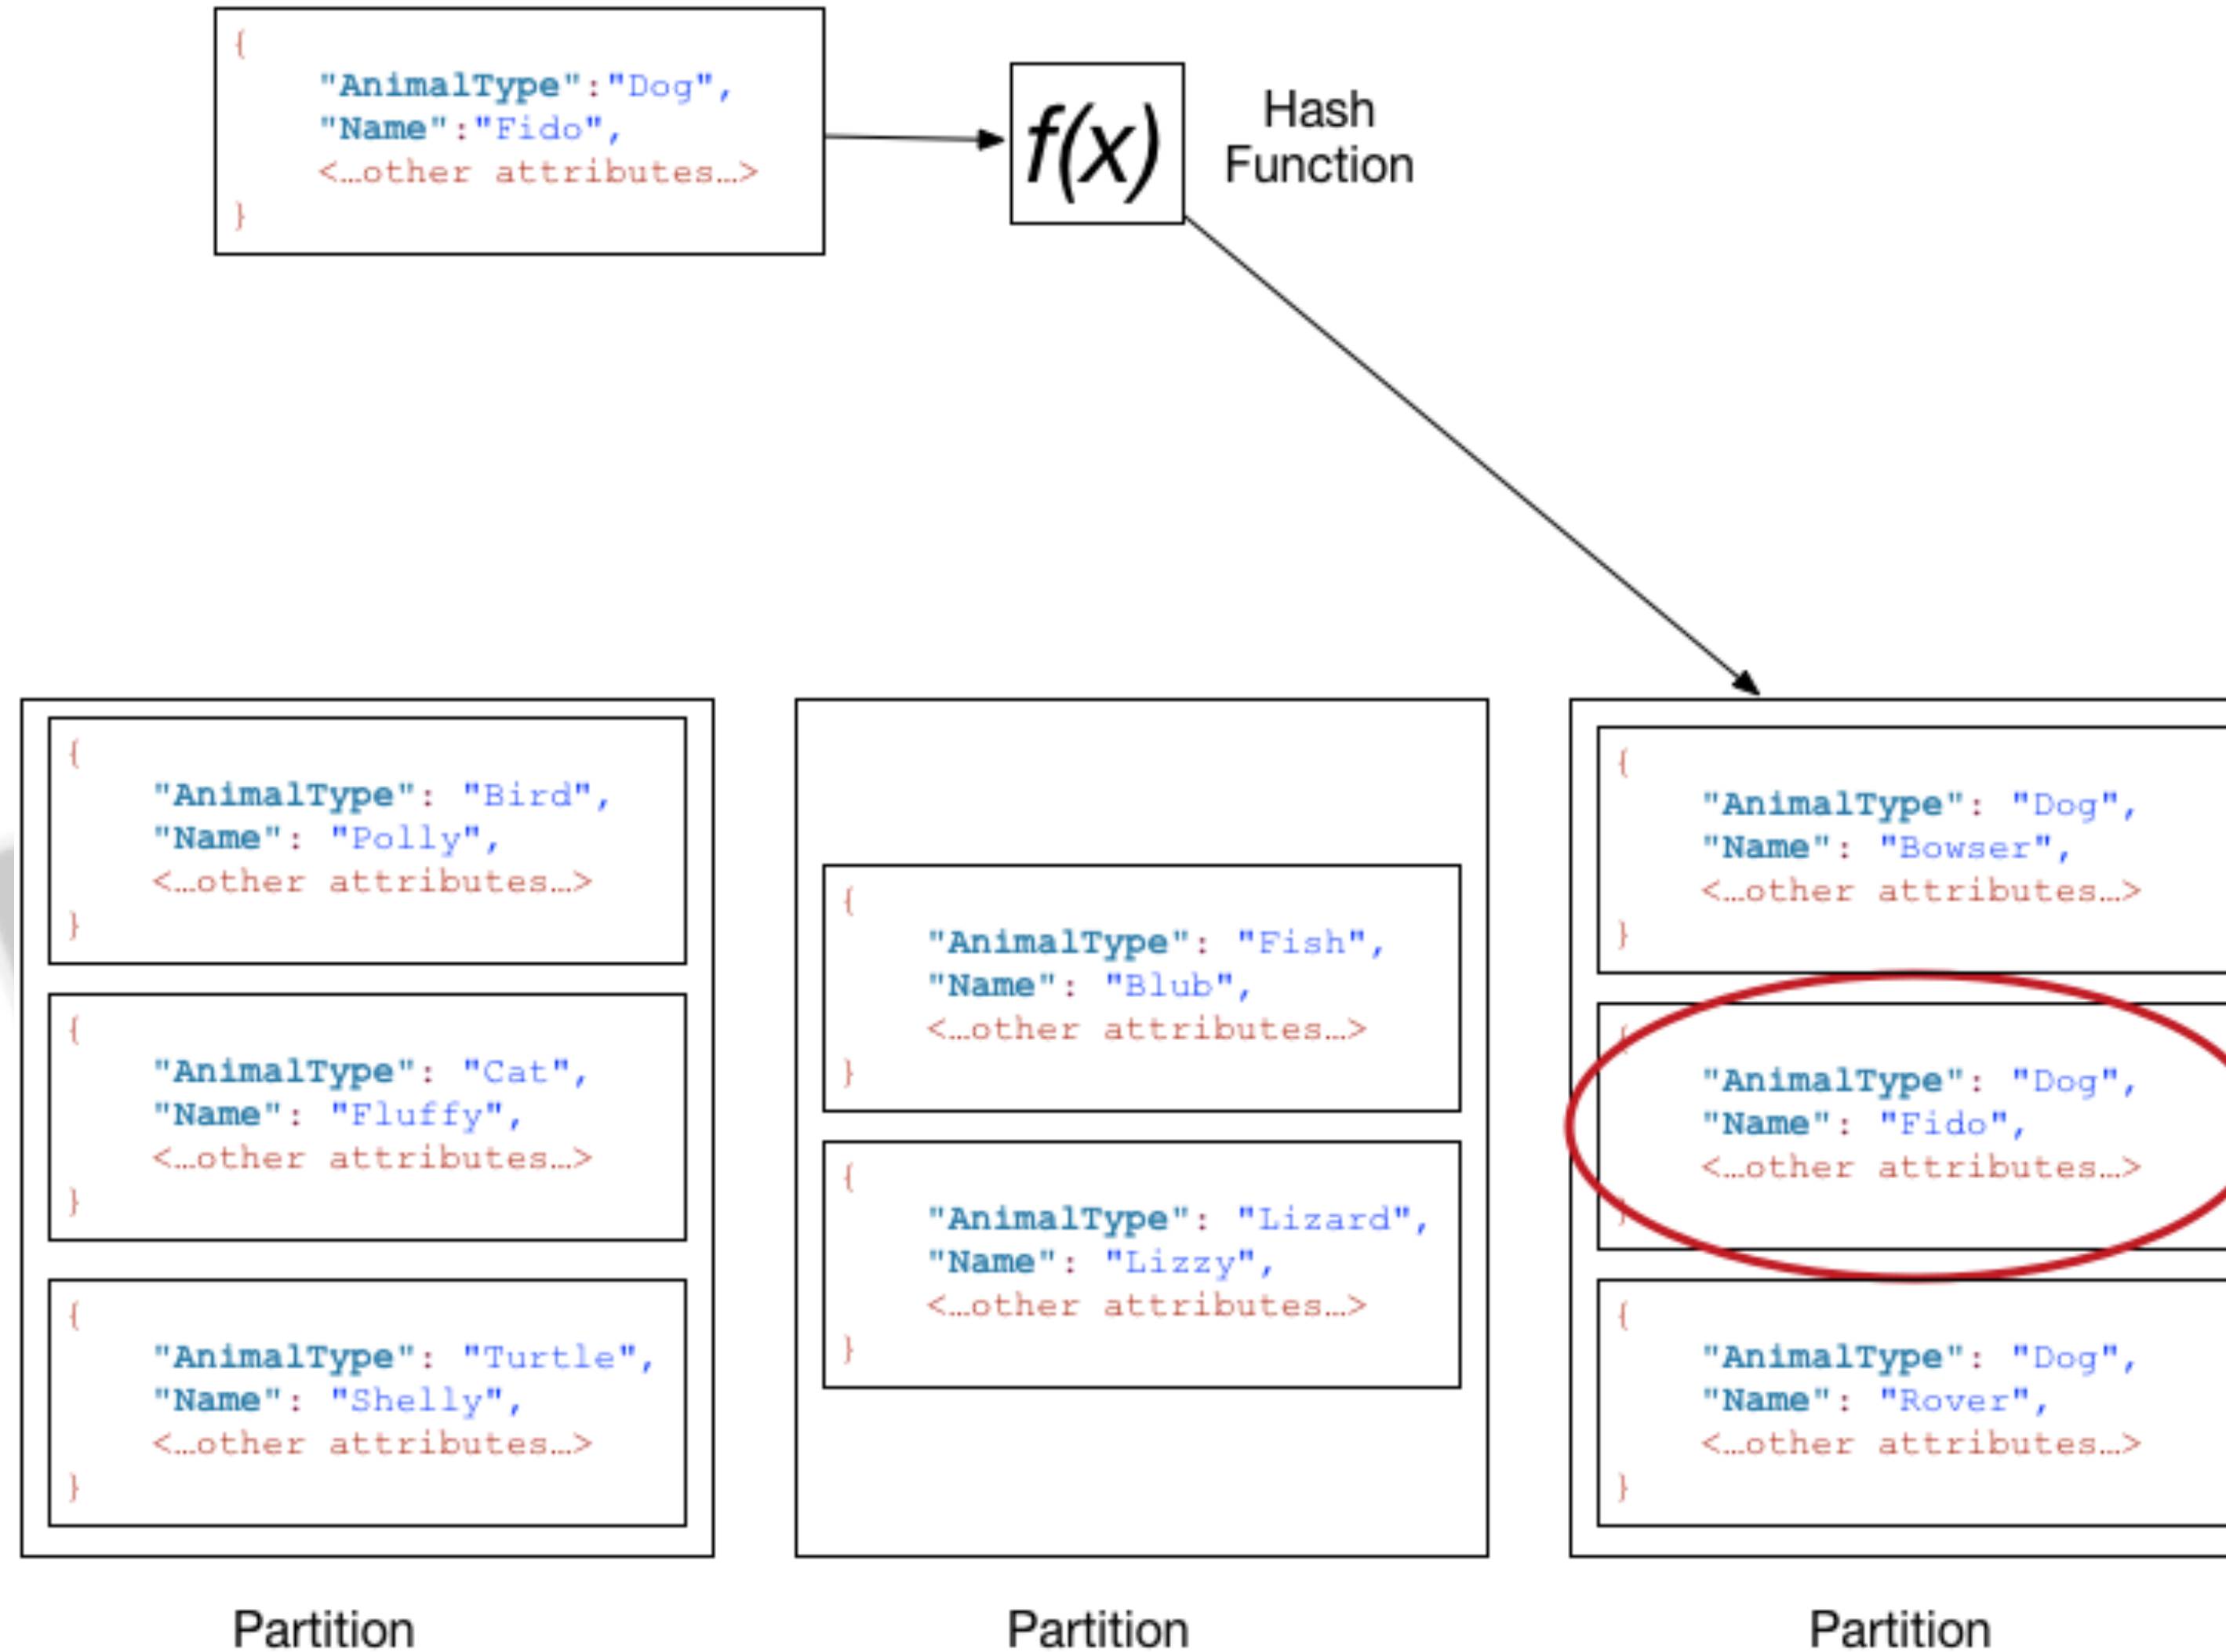
READ CONSISTENCY

- Tables in different regions are independent
- Eventually consistency reads (default)
 - When you read data from a DynamoDB table, the response might not reflect the results of a recently completed write operation. The response might include some stale data. If you repeat your read request after a short time, the response should return the latest data.
- Strongly consistent reads
 - When you request a strongly consistent read, DynamoDB returns a response with the most up-to-date data, reflecting the updates from all prior write operations that were successful. A strongly consistent read might not be available in the case of a network delay or outage.

PROVISIONED THROUGHPUT

- Read Unit
 - 1 read unit = one strong consistency read per second or two eventually consistency reads per second, up to 4KB
 - $3 \text{ KB} / 4 \text{ KB} = 0.75 \rightarrow 1$
 $6 \text{ KB} / 4 \text{ KB} = 1.5 \rightarrow 2$
1 read capacity unit per item $\times 80$ reads per second = 80
2 read capacity units per item $\times 100$ reads per second = 200
- Write unit
 - 1 write unit = one write per second up to 1KB
 - $512 \text{ bytes} / 1 \text{ KB} = 0.5 \rightarrow 1$
 $1.5 \text{ KB} / 1 \text{ KB} = 1.5 \rightarrow 2$
1 write capacity unit per item $\times 100$ writes per second = 100
2 write capacity units per item $\times 10$ writes per second = 20

PARTITIONS



EXAMPLES: CREATE A TABLE

```
CREATE TABLE Music (  
    Artist VARCHAR(20) NOT NULL,  
    SongTitle VARCHAR(30) NOT NULL,  
    AlbumTitle VARCHAR(25),  
    Year INT,  
    Price FLOAT,  
    Genre VARCHAR(10),  
    Tags TEXT,  
    PRIMARY KEY(Artist, SongTitle)  
);
```

```
{  
    TableName : "Music",  
    KeySchema: [  
        {  
            AttributeName: "Artist",  
            KeyType: "HASH", //Partition key  
        },  
        {  
            AttributeName: "SongTitle",  
            KeyType: "RANGE" //Sort key  
        }  
    ],  
    AttributeDefinitions: [  
        {  
            AttributeName: "Artist",  
            AttributeType: "S"  
        },  
        {  
            AttributeName: "SongTitle",  
            AttributeType: "S"  
        }  
    ],  
    ProvisionedThroughput: {  
        ReadCapacityUnits: 1,  
        WriteCapacityUnits: 1  
    }  
}
```

EXAMPLE: READ AN ITEM

```
SELECT * FROM Music WHERE Artist='No One You Know' AND SongTitle = 'Call Me Today'
```

```
{  
  TableName: "Music",  
  Key: {  
    "Artist": "No One You Know",  
    "SongTitle": "Call Me Today"  
  }  
}
```


EXAMPLE: QUERY A TABLE

```
SELECT * FROM Music
WHERE Artist='No One You Know' AND SongTitle = 'Call Me Today';

{
  TableName: "Music",
  KeyConditionExpression: "Artist = :a and SongTitle = :t",
  ExpressionAttributeValues: {
    ":a": "No One You Know",
    ":t": "Call Me Today"
  }
}
```

EXAMPLE: SCAN A TABLE

```
SELECT Artist, Title FROM Music;
```

```
{
```

```
  TableName: "Music",
```

```
  ProjectionExpression: "Artist, Title"
```

```
}
```

EXAMPLE: SCAN A TABLE

```
UPDATE Music
```

```
SET RecordLabel = 'Global Records'
```

```
WHERE Artist = 'No One You Know' AND SongTitle = 'Call Me Today';
```

```
{
  TableName: "Music",
  Key: {
    "Artist": "No One You Know",
    "SongTitle": "Call Me Today"
  },
  UpdateExpression: "SET RecordLabel = :label",
  ExpressionAttributeValues: {
    ":label": "Global Records"
  }
}
```


EXAMPLE: SCAN A TABLE

```
DELETE FROM Music
```

```
WHERE Artist = 'The Acme Band' AND SongTitle = 'Look Out, World';
```

```
{  
  TableName: "Music",  
  Key: {  
    Artist: "The Acme Band",  
    SongTitle: "Look Out, World"  
  }  
}
```



BITTIGER

PROGRAMMING INTERFACES: LOW-LEVEL

POST / HTTP/1.1

Host: dynamodb.<region>.<domain>;

Accept-Encoding: identity

Content-Length: <PayloadSizeBytes>

User-Agent: <UserAgentString>

Content-Type: application/x-amz-json-1.0

Authorization: AWS4-HMAC-SHA256 Credential=<Credential>, SignedHeaders=<Headers>, Signature=<Signature>

X-Amz-Date: <Date>

X-Amz-Target: DynamoDB_20120810.GetItem

```
{  
  "TableName": "Pets",  
  "Key": {  
    "AnimalType": {"S": "Dog"},  
    "Name": {"S": "Fido"}  
  }  
}
```

PROGRAMMING INTERFACES: DOCUMENT AND OBJECT

- Java and .NET Only

```
public class MusicDocumentDemo {  
  
    public static void main(String[] args) {  
  
        AmazonDynamoDBClient client = \  
            new AmazonDynamoDBClient();  
        DynamoDB docClient = new DynamoDB(client);  
  
        Table table = docClient.getTable("Music");  
        GetItemOutcome outcome = table.getItemOutcome(  
            "Artist", "No One You Know",  
            "SongTitle", "Call Me Today");  
  
        int year = outcome.getItem().getInt("Year");  
  
    }  
}
```

```
@DynamoDBTable(tableName="ProductCatalog")  
public class CatalogItem {  
  
    private Integer id;  
    private String title;  
    private String ISBN;  
    private Set<String> bookAuthors;  
    private String someProp;  
  
    @DynamoDBHashKey(attributeName="Id")  
    public Integer getId() { return id;}  
    public void setId(Integer id) {this.id = id;}  
  
    @DynamoDBAttribute(attributeName="Title")  
    public String getTitle() {return title; }  
    public void setTitle(String title) { this.title = title; }  
}
```


“

Whenever you find yourself on the side of majority, it is time to pause and reflect.

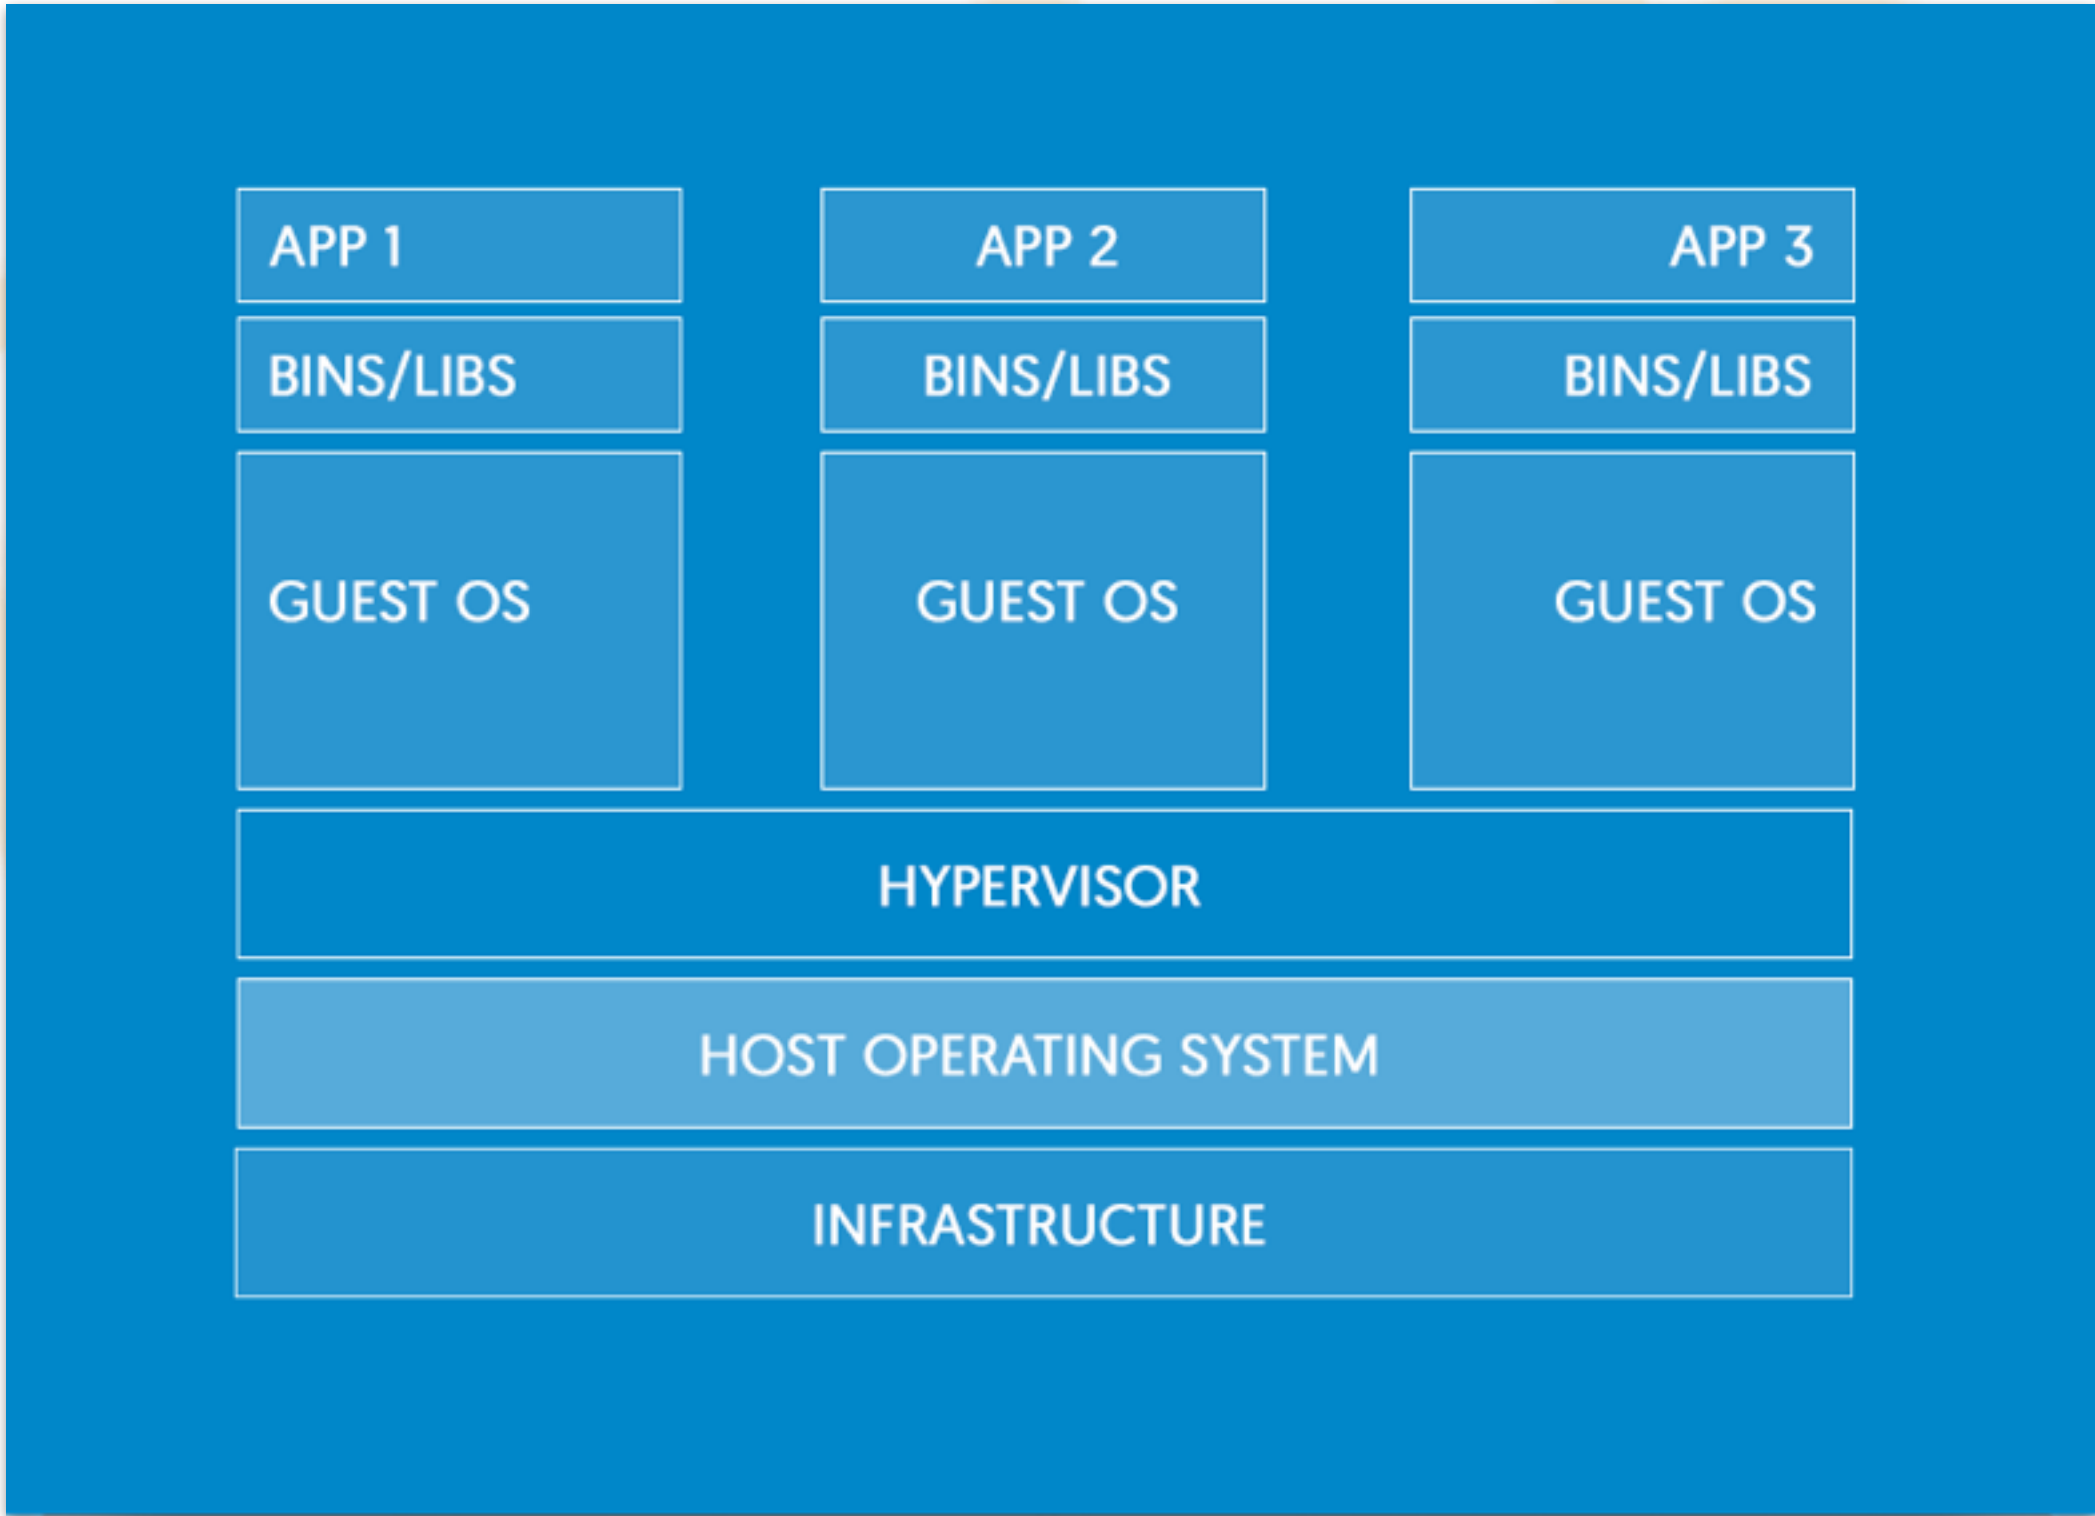
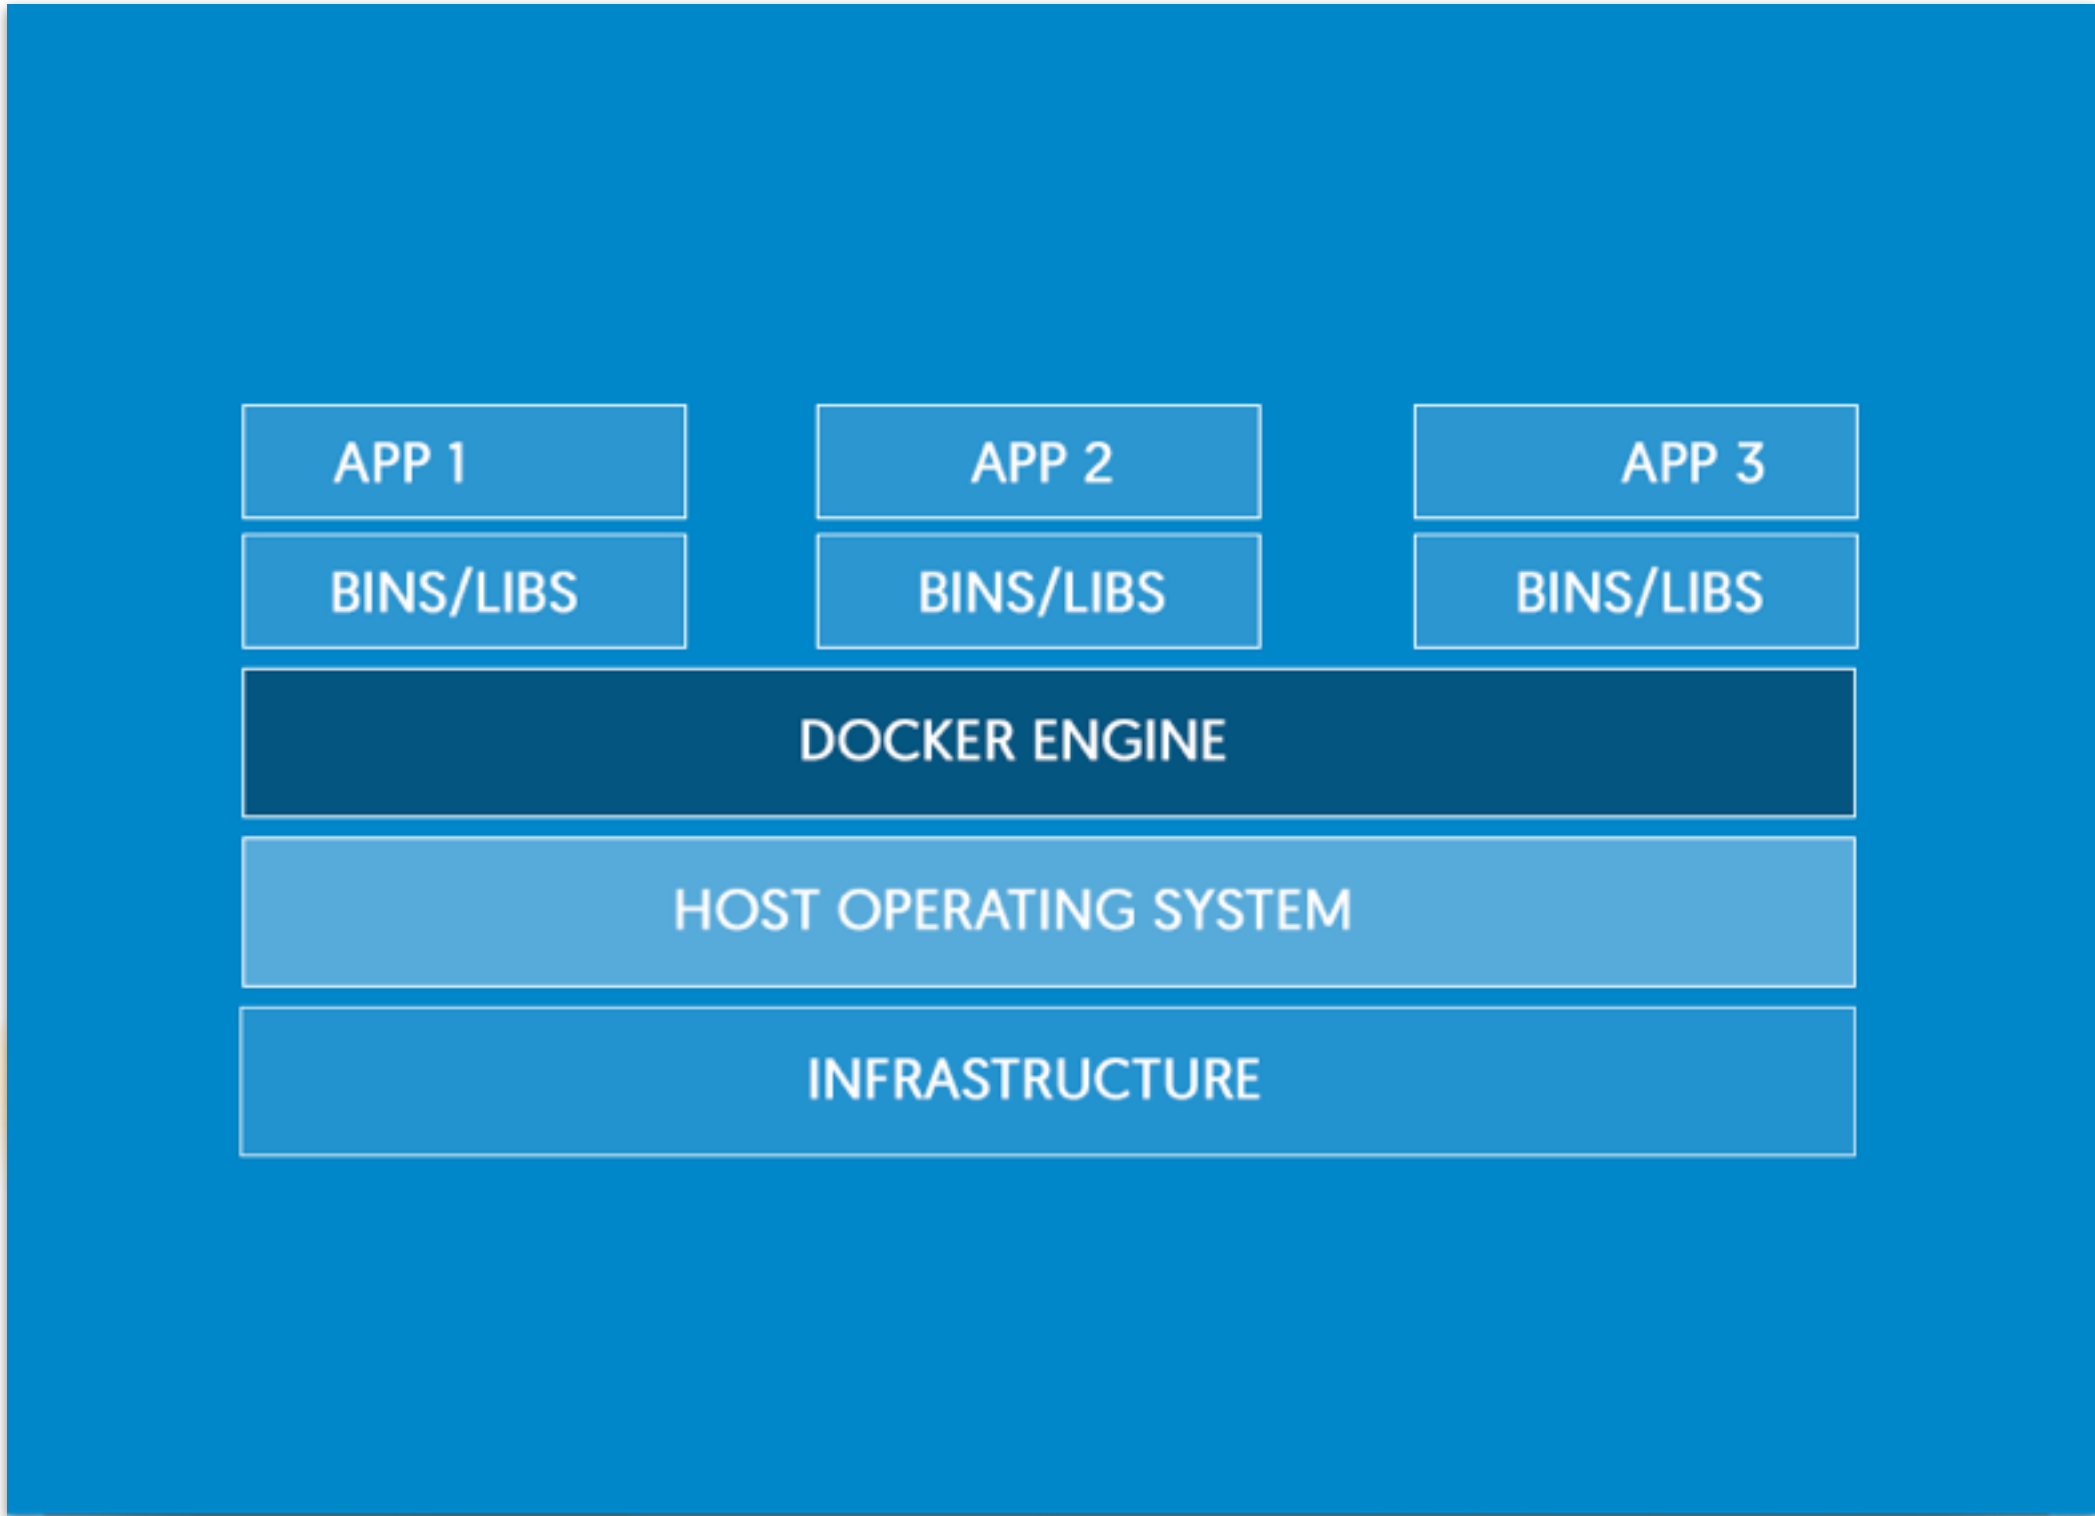
-Mark Twain



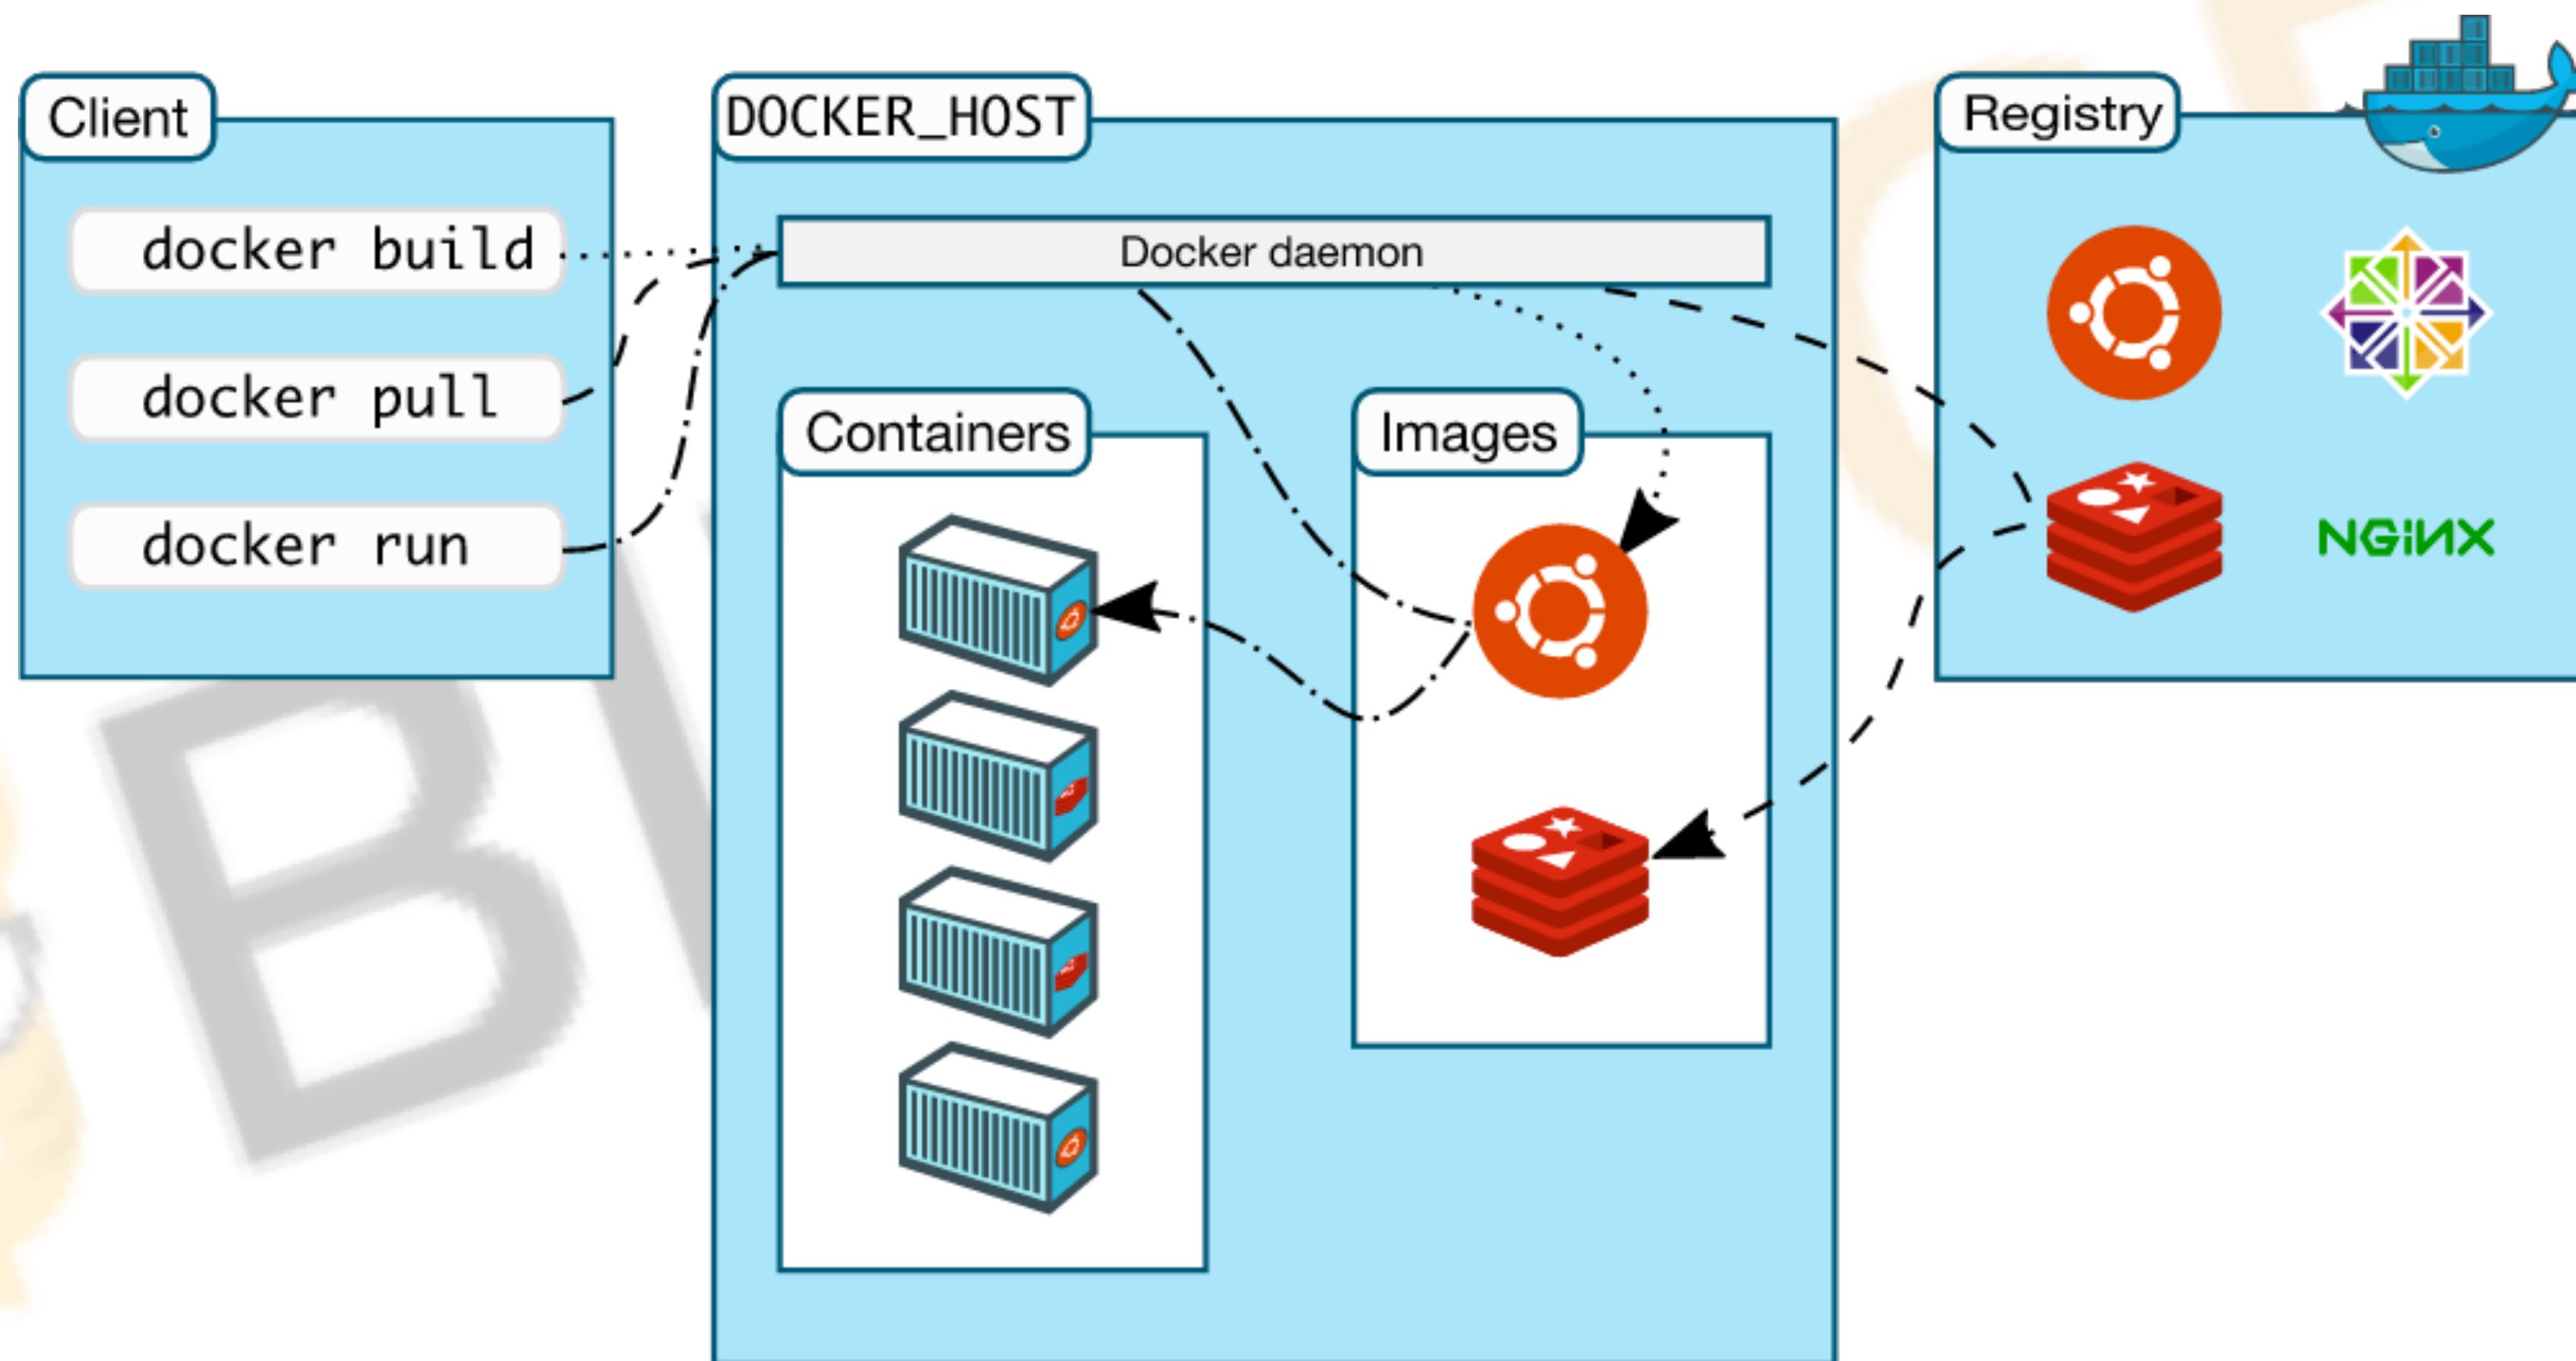
BITTIGER

DOCKER

DOCKER VS. VM



ARCHITECTURE



WHY MICROSERVICE? I MEAN, FUNDAMENTALLY

- Think about functions and process pool in programming language
 - Easy to write, debug, reuse a function
 - Functions talk to each other by parameters and return values
- Advantages
 - Isolation
 - Scalability, of course for deployment, not performance
- Disadvantages
 - Operational overhead
 - Performance overhead
 - I/O, error handling
 - Compatibility

ENABLE TECHNOLOGY

- Process separation, not hardware separation
- Namespace feature provided by Linux kernel
 - pid: Process isolation
 - net: network interface
 - ipc: interprocess communication
 - mnt: filesystem mount points
 - uts: kernel and version identifiers
- Control groups
- Union File Systems
- Container Format

KEY CONCEPTS

- Docker image
 - Like AMI + Packer
- Docker engine
 - Like Ansible
- Docker hub and registry
 - Like github



BITTIGER

DOCKERFILE

- FROM
- RUN
- CMD
- ENTRYPOINT
- ADD/COPY
- EXPOSE
- .dockerignore

FROM ubuntu

Install vnc, xvfb in order to create a 'fake' display and firefox

RUN apt-get update && apt-get install -y x11vnc xvfb firefox

RUN mkdir ~/.vnc

Setup a password

RUN x11vnc -storepasswd 1234 ~/.vnc/passwd

Autostart firefox (might not be the best way, but it does the trick)

RUN bash -c 'echo "firefox" >> ~/.bashrc'

EXPOSE 5900

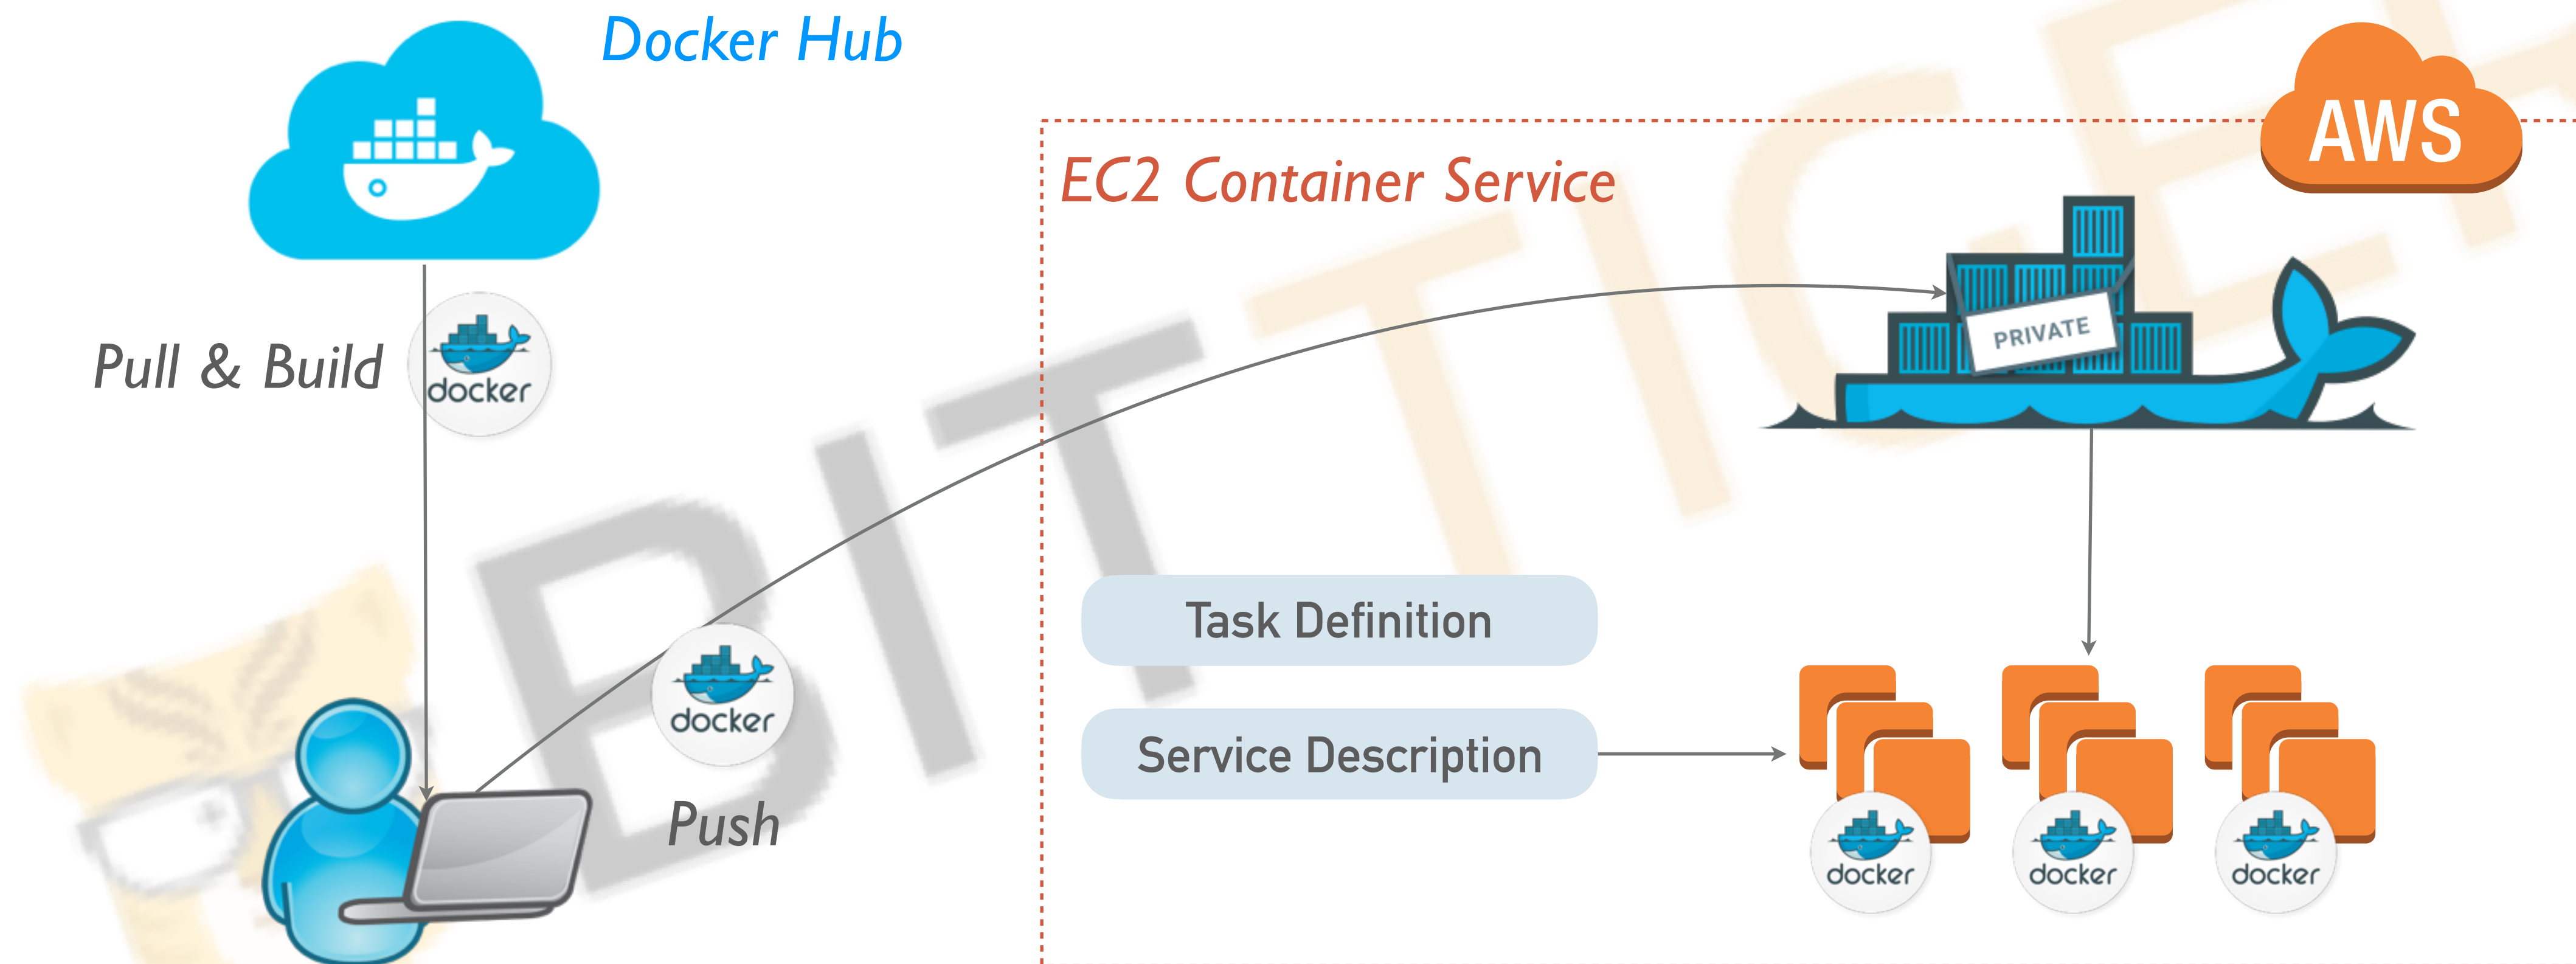
CMD ["x11vnc", "-forever", "-usepw", "-create"]



BITTIGER

ECS

DOCKER ON EC2



CONTAINER INSTANCE

- Using AMI with agent enabled
 - Amazon ECS-optimized AMI (http://docs.aws.amazon.com/AmazonECS/latest/developerguide/container_agent_versions.html#ecs-optimized-ami-agent-versions)
 - No more HVM and PV
- IAM Role: AmazonEC2ContainerServiceforEC2Role



TASK DEFINITION

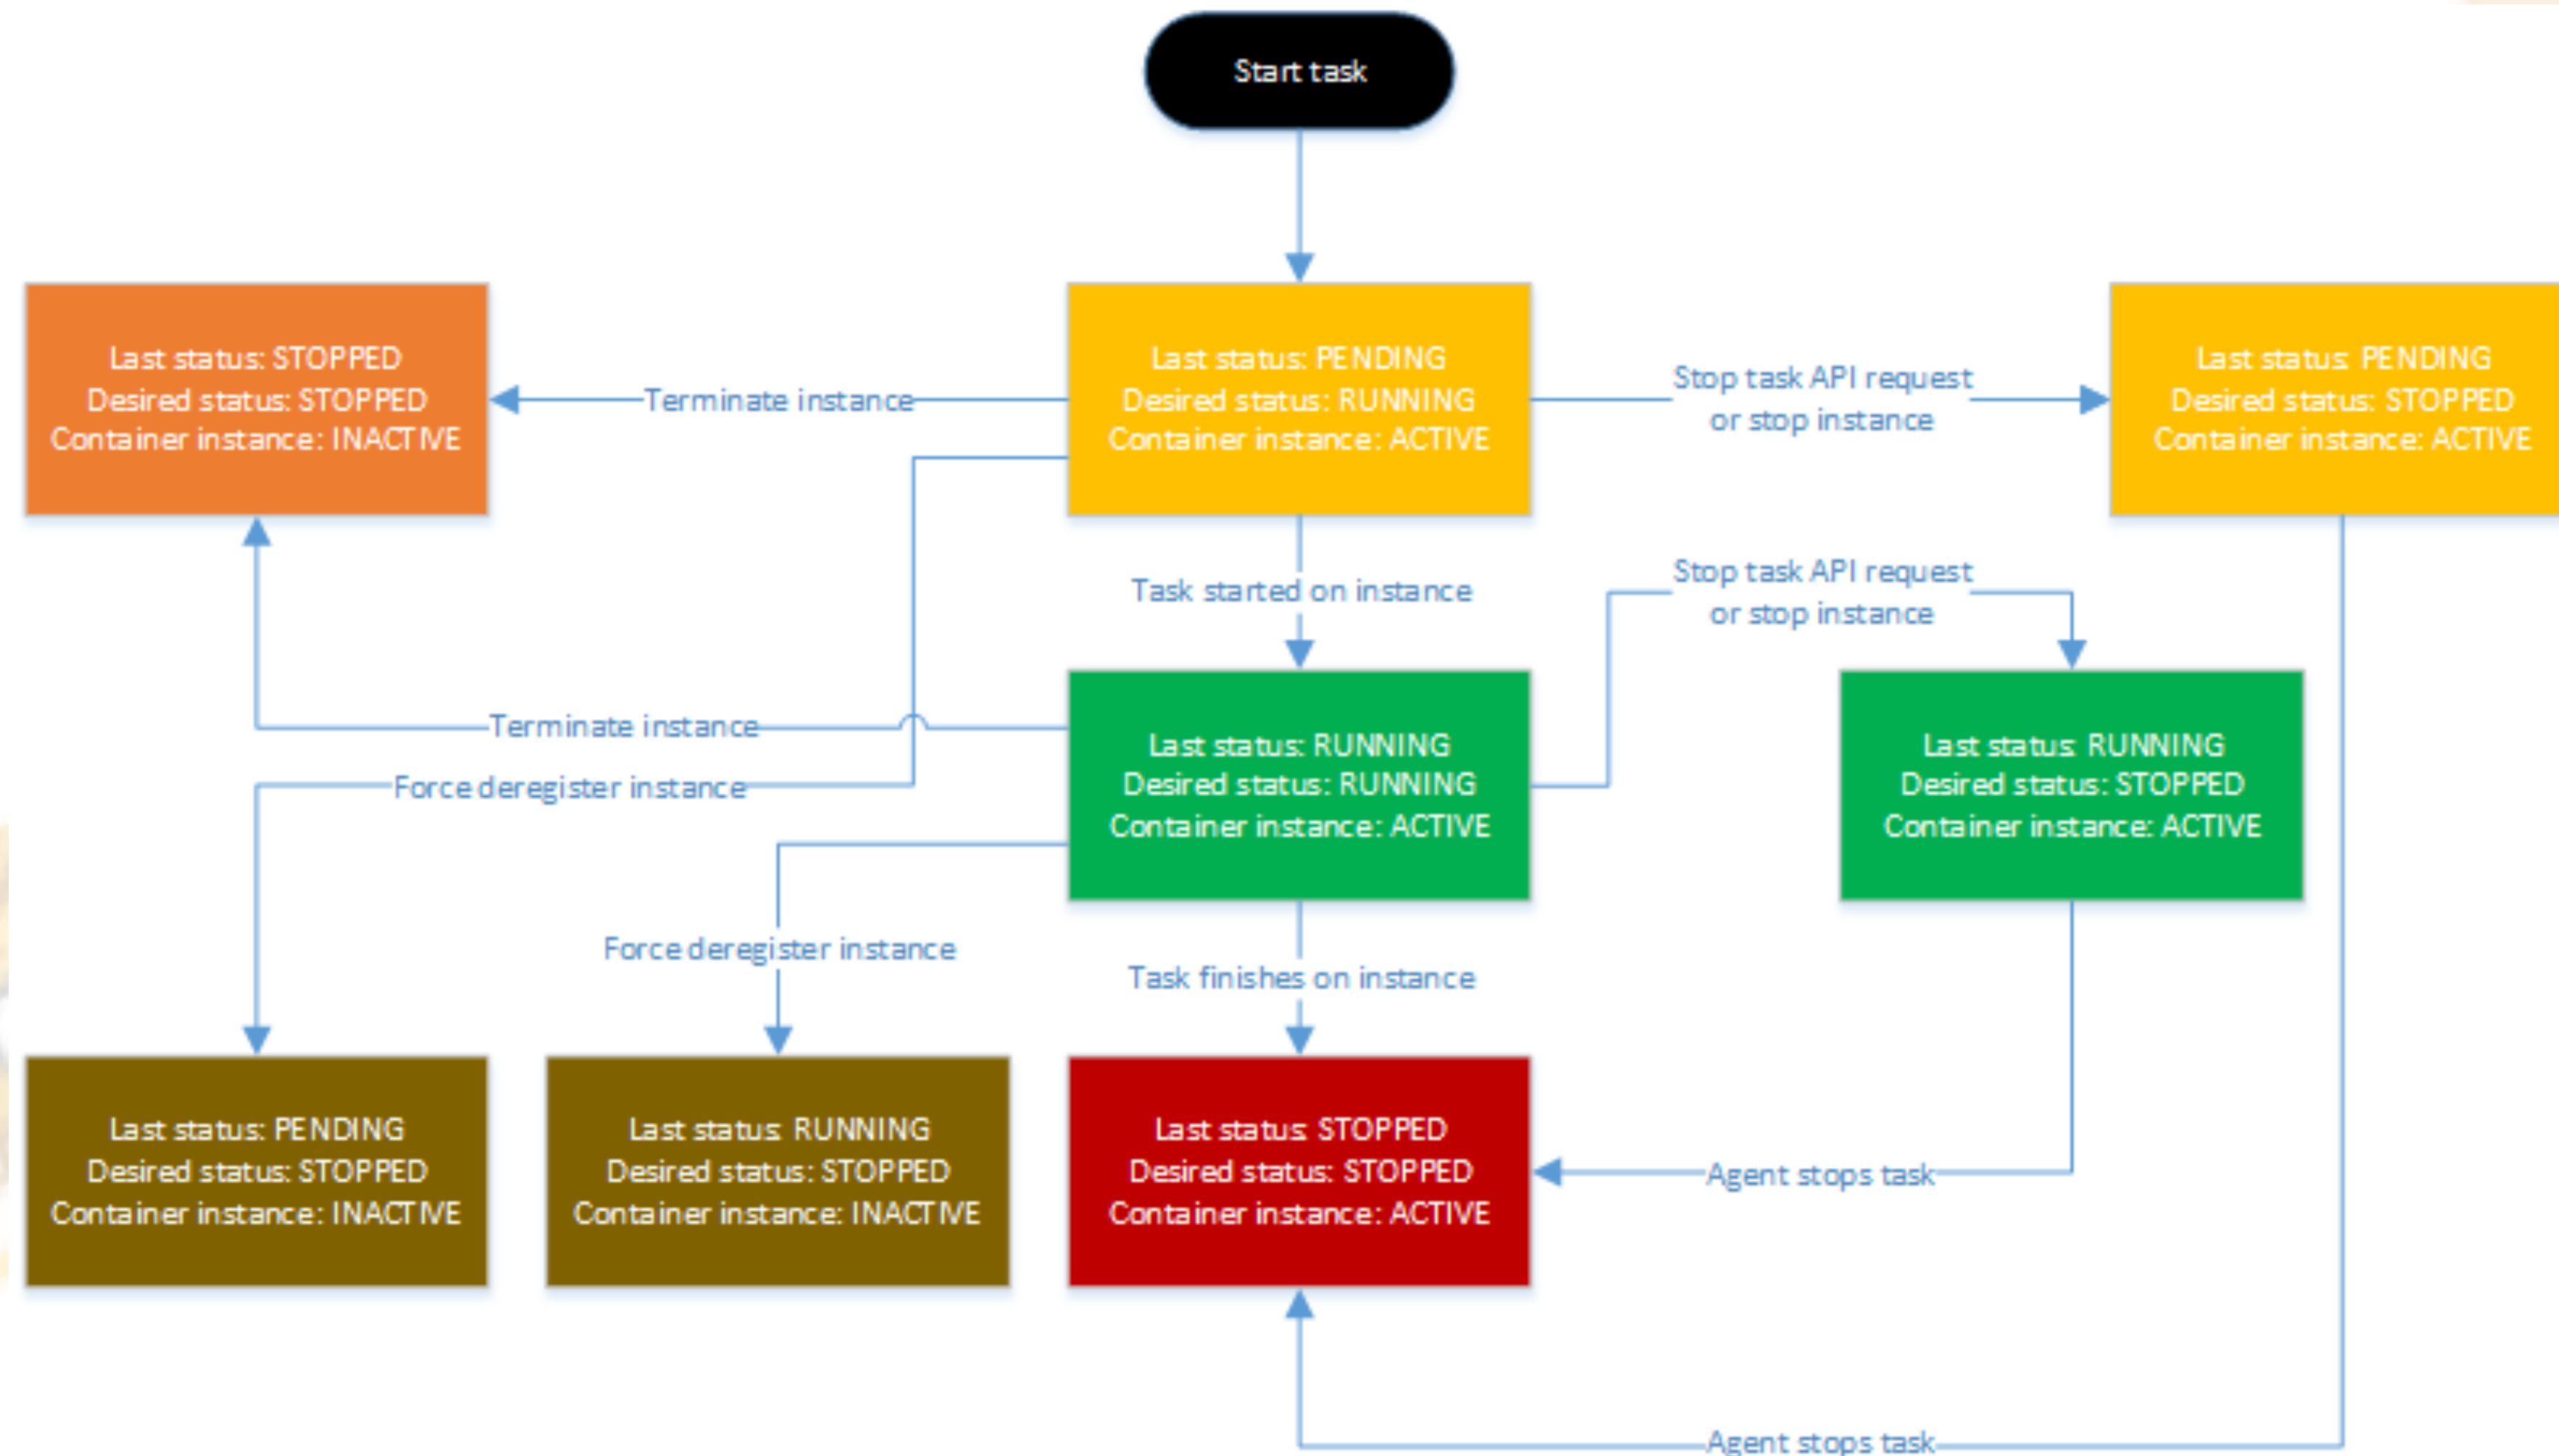
- Specification of run Docker containers

```
{
  "containerDefinitions": [
    {
      "name": "sample-app",
      "image": "123456789012.dkr.ecr.us-west-2.amazonaws.com/aws-nodejs-sample:v1",
      "memory": "200",
      "cpu": "10",
      "essential": true
    }
  ],
  "family": "example_task_3",
  "taskRoleArn": "arn:aws:iam::123456789012:role/AmazonECSTaskS3BucketRole"
}
```


TASK SCHEDULING


- Run task with scheduler
 - Specify how many tasks (workers) you need
- Task placement
 - binpack
 - Least available amount of CPU or memory of instance
 - random
 - spread
 - Based on specific attributed, such as availability zones

TASK LIFECYCLE



SERVICE

- A set of instance running given tasks
 - The number of tasks are maintained to desired level
 - Service can be behind a load balancer
 - Auto-balancing across availability zones
 - Stop service in a optimal way across availability zones
 - Stop instance where its zones as more instance running



NYC TAXI



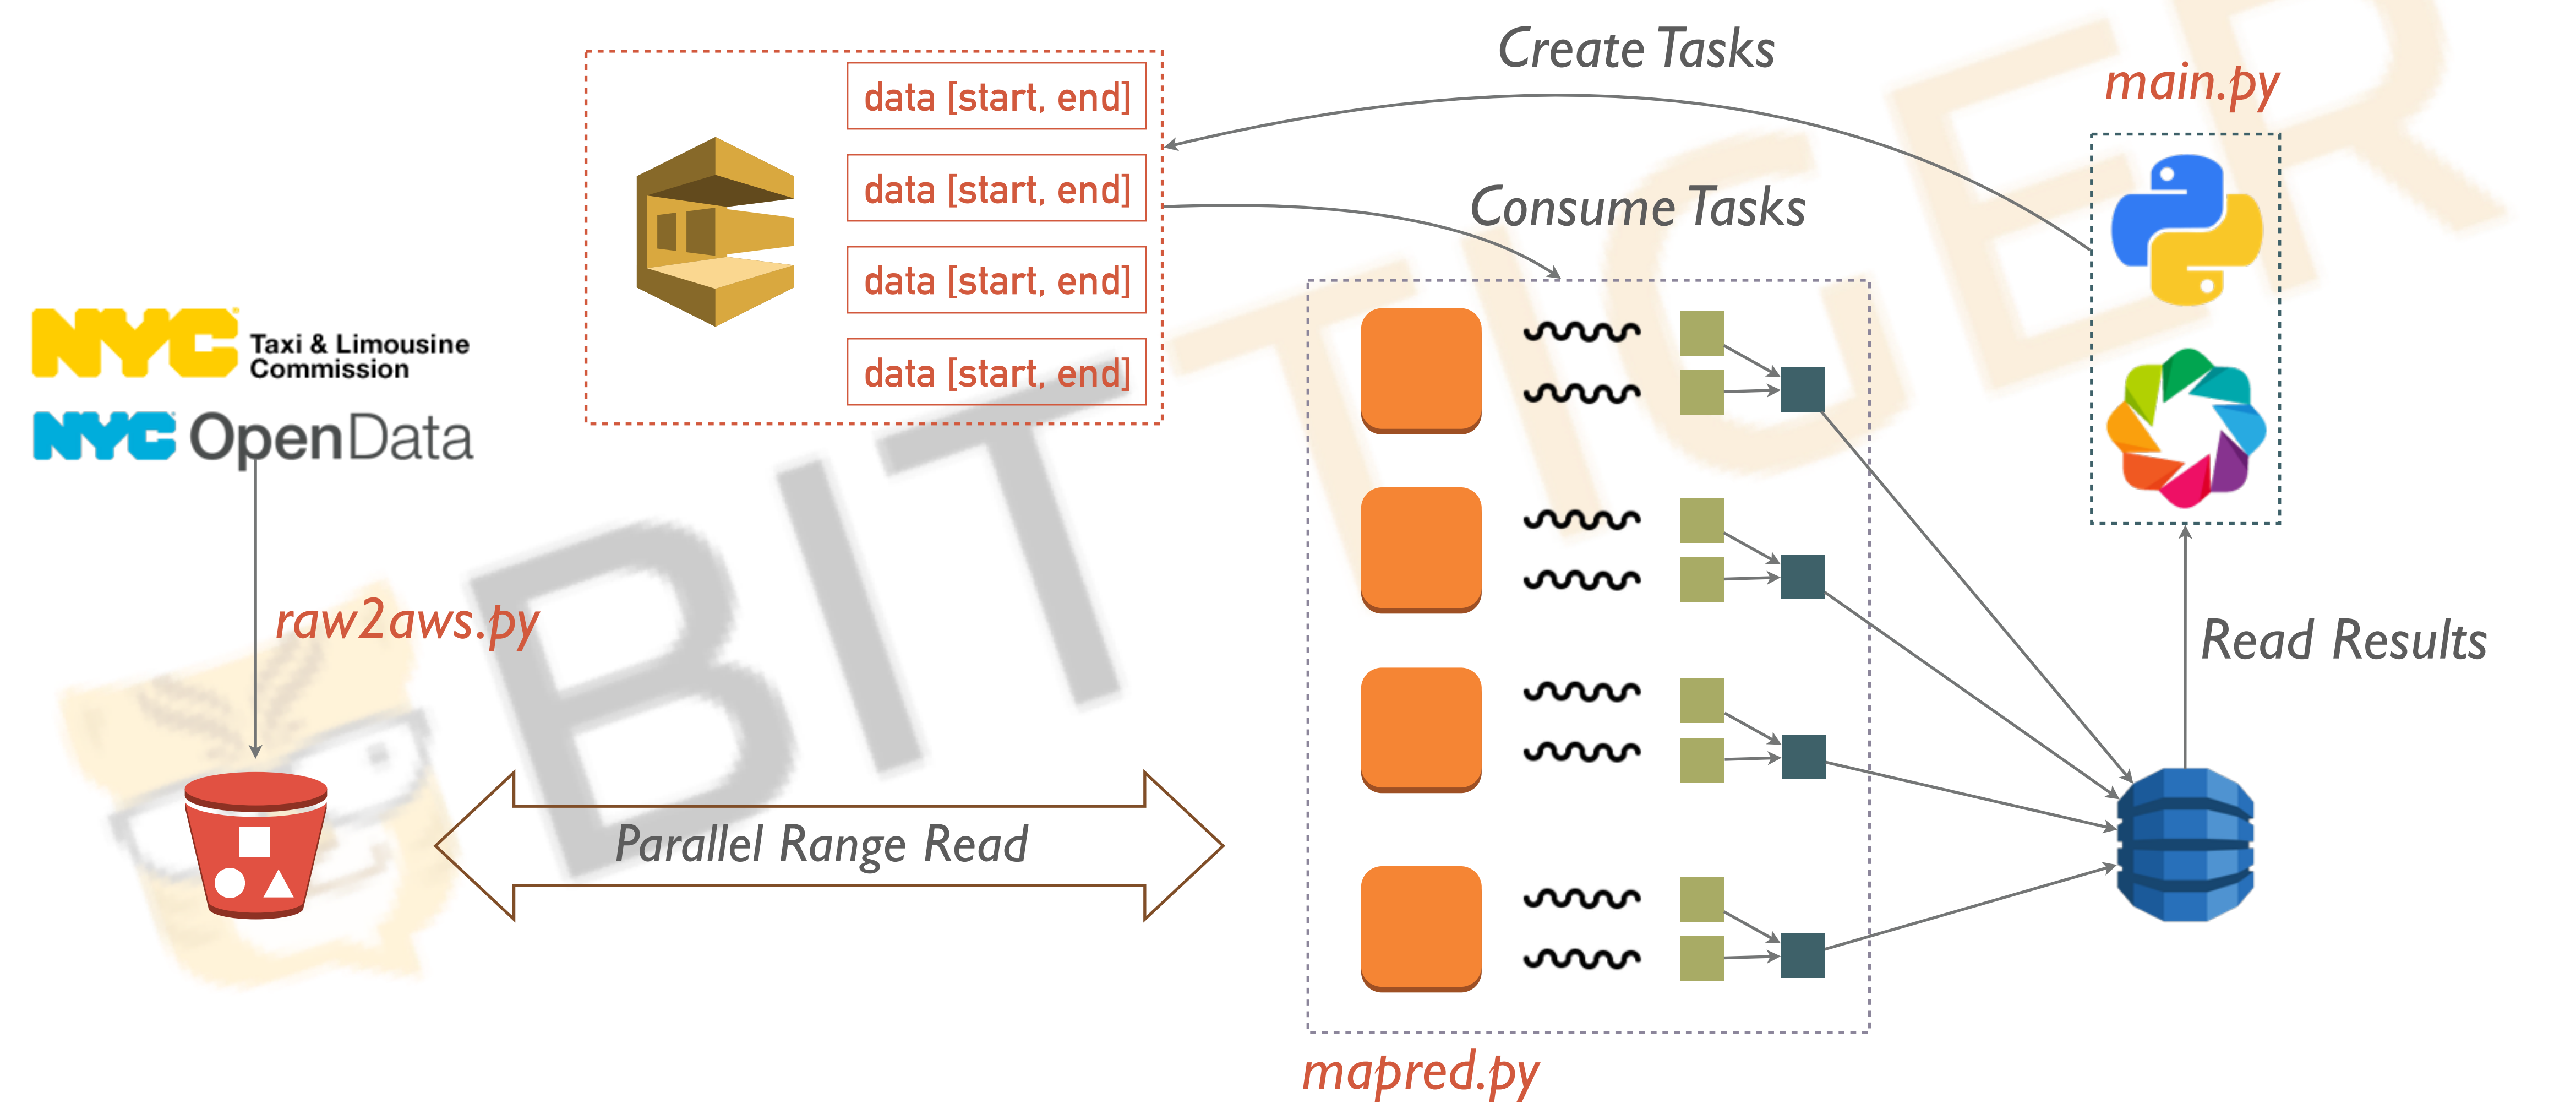
WHAT WE HAVE NOW

- Well-formate input data
 - 80 bytes per record
- A to map-reduce program to count statistics
 - Calculate each trip in which districts
- How to connect them together to process data in parallel?
 - A queue to provide concurrency and task tracking
 - A result table to aggregate statistics

MAJOR COMPUTATION

- Test is a point is in polygon
 - A point (longitude, latitude)
 - A polygon described by a series of boundary points
 - <http://erich.realtimerendering.com/ptinpoly/>
 - Python shapely package: `Polygon.contains(Point())`
- Geojson Data
 - <https://github.com/dwillis/nyc-maps>
- 99% computation

WORKFLOW

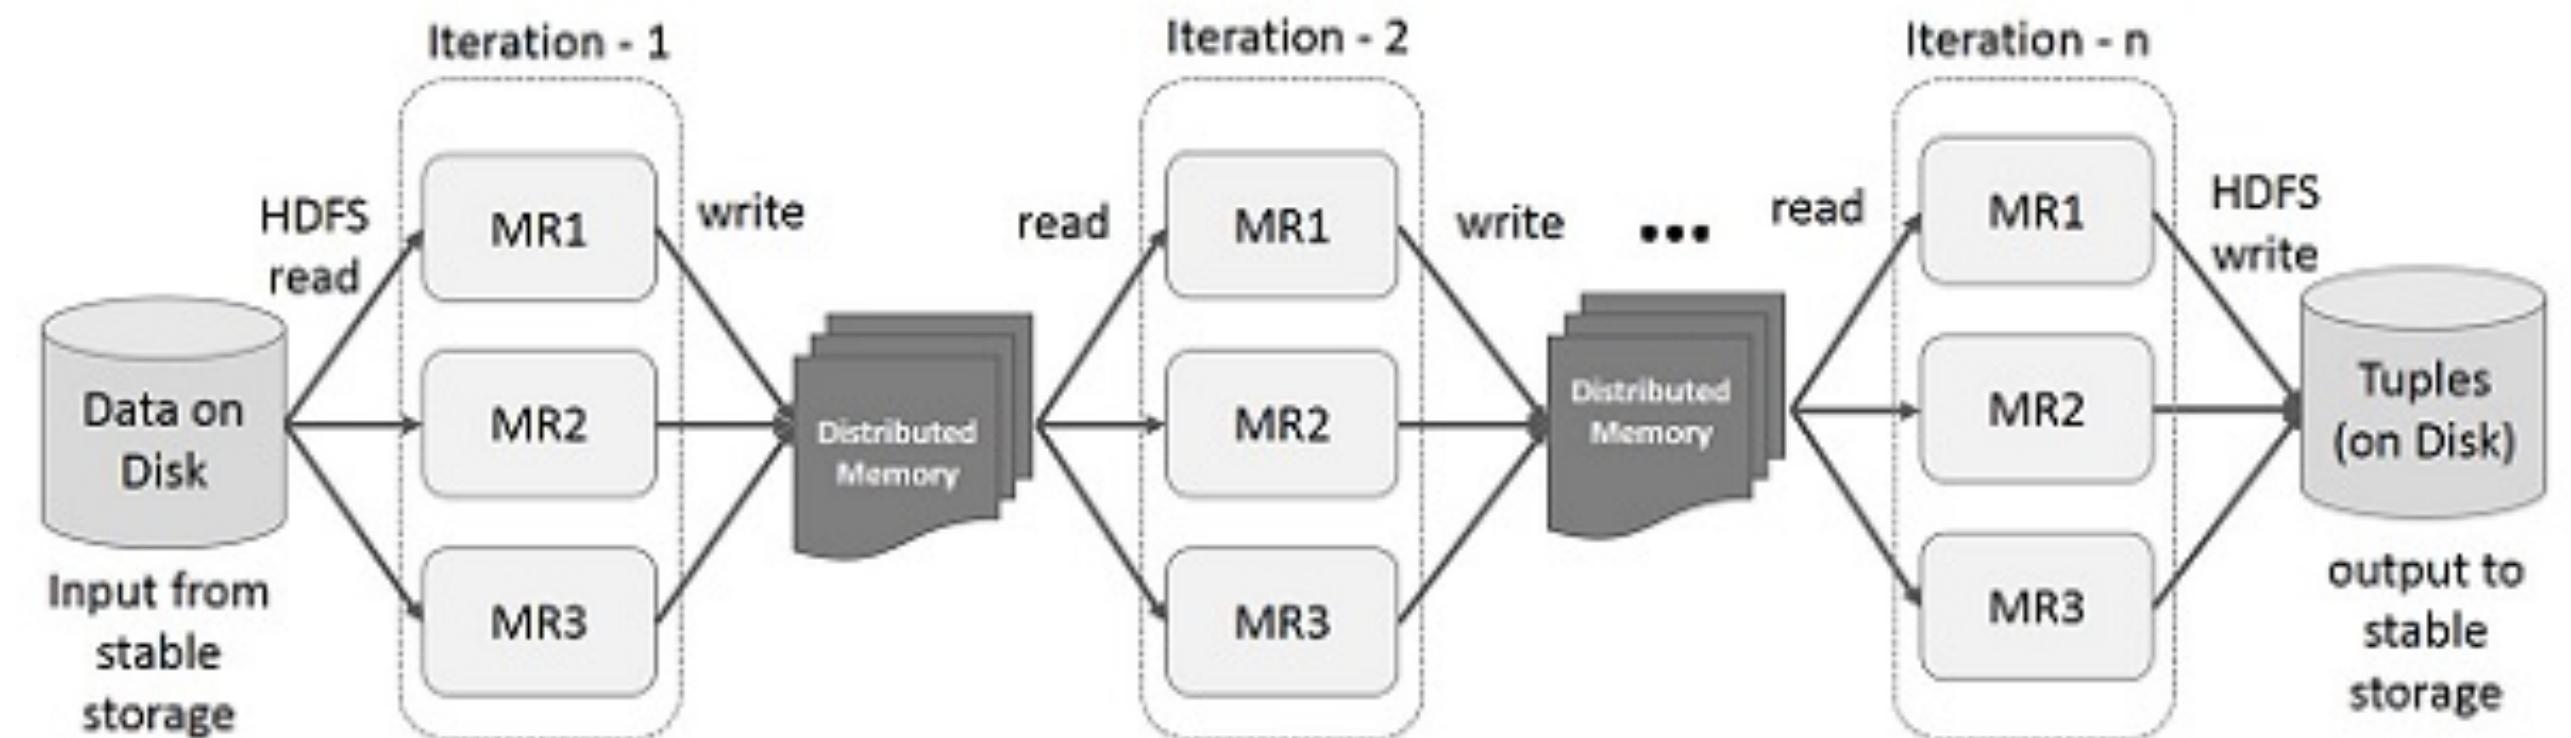
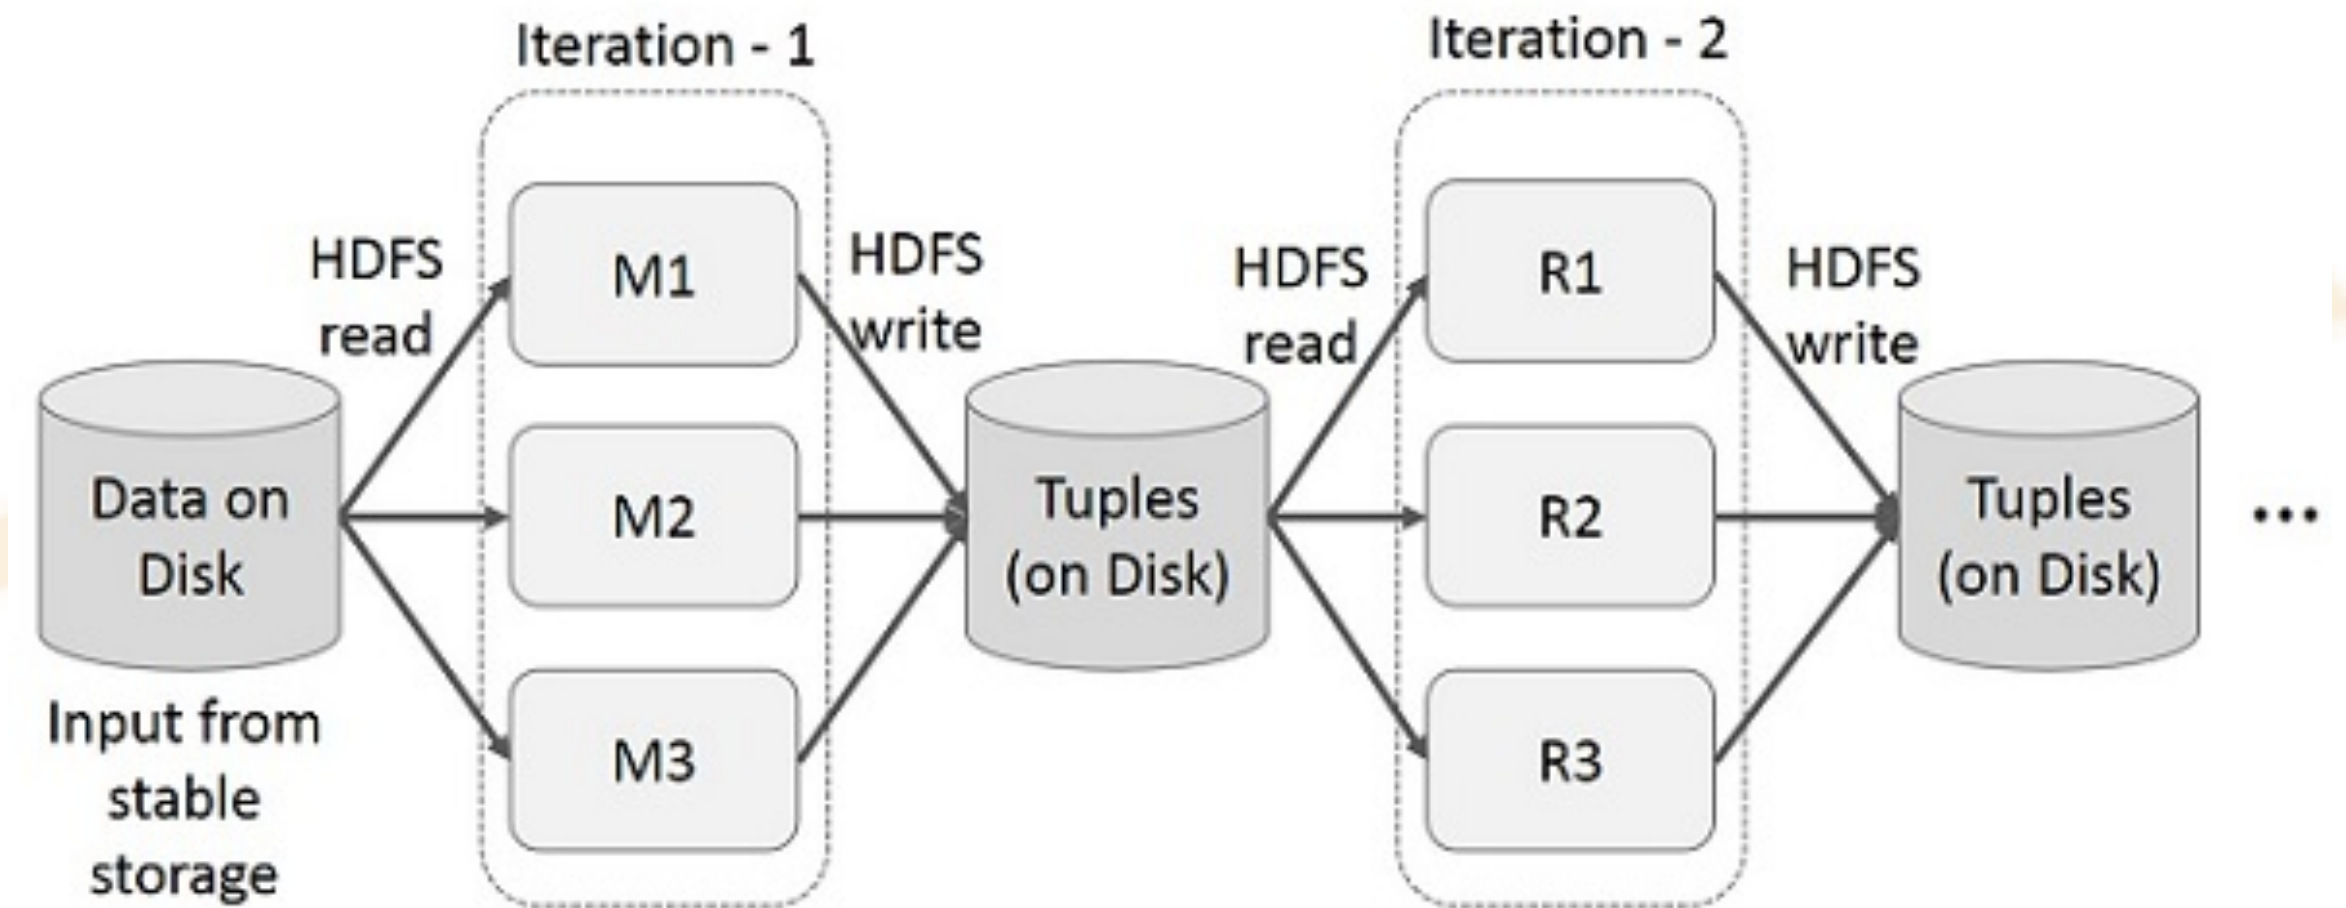


DEPLOYMENT

- Classic
 1. Build AMI
 2. Terraform for infrastructure
 3. Configure instances via user-data or Ansible
 4. Run tasks via user-data
- Container (in a infrastructure)
 1. Build Docker Image and push to registry
 2. Define tasks and service
 3. Run

WHY NOT SPARK?

- We have **no sharing state/data** at all!
 - Why bother using RDD?!
- We have only one step!
- Overhead to run on top of Spark
- Too complex for our tasks
- For this task, we are much faster than using Spark




USE CONTAINER?

- All tasks are homogenous
 - Load from one single executable
- All tasks runs at full speed of CPU
- We can run as many as processes if memory allows...
- Low-level process/thread management is more efficient
- Container not idea for this case? but VM overhead
 - Overhead of packing and deploy
 - Overhead of isolation
 - Overhead of Docker daemon

BOKEH ボケ

- A Python Interactive visualization library mainly for web
- Translate to Javascript
- Examples used in project
 - <https://demo.bokehplots.com/apps/movies>
 - <http://bokeh.pydata.org/en/latest/docs/gallery/texas.html>
 - <https://www.quantinsti.com/blog/python-data-visualization-using-bokeh/>



QUESTIONS

- bittiger-aws@googlegroups.com



BITTIGER

Copyright 2017, Nan Dun, all rights reserved